



OASIS3 User Guide

OASIS3-MCT_1.0

Edited by:

*S. Valcke, T. Craig, L. Coquart
CERFACS/CNRS SUC URA No1875*

CERFACS TR/CMGC/XX/XX

July 2012

Copyright Notice

© Copyright 2012 by CERFACS

All rights reserved.

No parts of this document should be either reproduced or commercially used without prior agreement by CERFACS representatives.

How to get assistance?

Assistance can be obtained as listed below.

Phone Numbers and Electronic Mail Addresses

Name	Phone	Affiliation	e-mail
Laure Coquart	+33-5-61-19-30-29	CNRS	oasishelp(at)cerfacs.fr
Sophie Valcke	+33-5-61-19-30-76	CERFACS	

How to get documentation ?

The documentation can be downloaded from the OASIS web site under the URL :

<http://oasis.enes.org>

Contents

0.1	Acknowledgments	1
1	Introduction	2
1.1	Step-by-step use of OASIS3-MCT	3
1.2	OASIS3-MCT sources	3
1.3	Licenses and Copyrights	3
1.3.1	OASIS3-MCT license and copyright statement	3
1.3.2	MCT copyright statement	4
1.3.3	The SCRIP 1.4 license copyright statement	5
2	Interfacing a component model with the OASIS3-MCT library	6
2.1	Use of OASIS3-MCT library	7
2.2	Initialisation	7
2.2.1	Coupling initialisation	7
2.2.2	Communicator for internal parallelisation	7
2.2.3	Coupling through a subset of the component model processes	8
2.3	Grid data file definition	9
2.4	Partition definition	11
2.4.1	Serial (no partition)	11
2.4.2	Apple partition	11
2.4.3	Box partition	11
2.4.4	Orange partition	14
2.5	Coupling field declaration	14
2.6	End of definition phase	15
2.7	Sending and receiving actions	15
2.7.1	Sending a coupling (or I/O) field	15
2.7.2	Receiving a coupling (or I/O) field	16
2.8	Termination	17
2.9	Auxiliary routines	17
2.10	Coupling algorithms - SEQ and LAG concepts	18
2.10.1	The lag concept	18
2.10.2	The sequence concept	21
3	The configuration file <i>namcouple</i>	22
3.1	An example of a simple <i>namcouple</i>	22
3.2	First section of <i>namcouple</i> file	24
3.3	Second section of <i>namcouple</i> file	25
3.3.1	Second section of <i>namcouple</i> for EXPORTED and EXPOUT fields	25
3.3.2	Second section of <i>namcouple</i> for OUTPUT fields	26
3.3.3	Second section of <i>namcouple</i> for INPUT fields	26
4	Transformations and interpolations	28

4.1	Time transformations	28
4.2	The pre-processing transformations	29
4.3	The interpolation	29
4.4	The “cooking” stage	33
4.5	The post-processing	34
5	OASIS3 auxiliary data files	35
5.1	Field names and units	35
5.2	Grid data files	35
5.3	Coupling restart files	36
5.4	Mapping files	36
5.5	Input data files	37
5.6	Transformation auxiliary data files	37
5.6.1	Auxiliary data files for SCRIPR	37
6	Compiling and running OASIS3 and TOYOASIS3	39
6.1	Compiling OASIS3 and debugging	39
6.1.1	Compilation with TopMakefileOasis3	39
6.1.2	CPP keys	39
6.1.3	Debugging	40
6.2	Running OASIS3 in coupled mode with TOYOASIS3	40
6.2.1	TOYOASIS3 description	40
6.2.2	Compiling and Running TOYOASIS3	42
6.3	Known problems when compiling or running OASIS3 on specific platforms	44
A	The grid types for the transformations	45
B	Changes between versions	46
B.1	Changes between oasis3_3 and oasis3-mct	46
B.2	Changes between oasis3_3 and oasis3_prism_2_5	48
B.3	Changes between oasis3_prism_2_5 and oasis3_prism_2_4	51
B.4	Changes between oasis3_prism_2_4 and oasis3_prism_2_3	52
B.5	Changes between oasis3_prism_2_3 and oasis3_prism_2_2	53
B.6	Changes between oasis3_prism_2_2 and oasis3_prism_2_1	53
B.7	Changes between oasis3_prism_2_1 and oasis3_prism_1_2	54
C	Coupled models realized with OASIS	56

0.1 Acknowledgments

We would like to thank the main past or present developers of OASIS are (in alphabetical order, with the name of their institution at the time of their contribution to OASIS):

Arnaud Caubel (LSCE/IPSL & FECIT/Fujitsu), Damien Declat (CERFACS), Italo Epicoco (CMCC), Veronika Gayler (MPI-M&D), Josefine Ghattas (CERFACS), Jean Latour (CERFACS & Fujitsu-Fecit), Eric Maisonnave (CERFACS), Silvia Mocavero (CMCC), Elodie Rapaport (CERFACS), Hubert Ritzdorf (CCRLE-NEC), Sami Saarinen (ECMWF), Eric Sevault (Météo-France), Laurent Terray (CERFACS), Olivier Thual (CERFACS), Reiner Vogelsang (SGI Germany), Li Yan (CERFACS).

We also would like to thank the following people for their help and suggestions in the design of the OASIS software (in alphabetical order, with the name of their institution at the time of their contribution to OASIS):

Dominique Astruc (IMFT), Chandan Basu (NSC, Sweden), Sophie Belamari (Météo-France), Dominique Bielli (Météo-France), Gilles Bourhis (IDRIS), Pascale Braconnot (IPSL/LSCE), Sandro Calmanti (Météo-France), Christophe Cassou (CERFACS), Yves Chartier (RPN), Jalel Chergui (IDRIS), Philippe Courtier (Météo-France), Philippe Dandin (Météo-France), Michel Déqué (Météo-France), Ralph Doescher (SMHI), Jean-Louis Dufresne (LMD), Jean-Marie Epitalon (CERFACS), Laurent Fairhead (LMD), Uwe Fladrich (SMHI), Marie-Alice Foujols (IPSL), Gilles Garric (CERFACS), Eric Guilyardi (CERFACS), Charles Henriët (CRAY France), Pierre Herchuelz (ACCRI), Maurice Imbard (Météo-France), Luis Kornblueh (MPI-M), Stephanie Legutke (MPI-M&D), Claire Lévy (LODYC), Olivier Marti (IPSL/LSCE), Sébastien Masson (IPSL/LOCEAN), Claude Mercier (IDRIS), Pascale Noyret (EDF), Andrea Piacentini (CERFACS), Marc Pontaud (Météo-France), Adam Ralph (ICHEC), René Redler (MPI-M), Tim Stockdale (ECMWF), Rowan Sutton (UGAMP), Véronique Taverne (CERFACS), Jean-Christophe Thil (UKMO), Nils Wedi (ECMWF).

Chapter 1

Introduction

In 1991, CERFACS started the development of a software interface to couple existing numerical General Circulation Models of the ocean and of the atmosphere. Today, the OASIS3.3 coupler, which is the result of more than 20 years of evolution is used by about 30 modelling groups in Europe, Australia, Asia and North America on the different computing platforms. The list of coupled models realized with OASIS3 and OASIS2 can be found in tables C.1 and C.2 in Appendix C.

OASIS sustained development is ensured by a collaboration between CERFACS and the Centre National de la Recherche Scientifique (CNRS) and its maintenance and user support is presently reinforced with additional resources coming from IS-ENES project funded by the EU (FP7 - GA no 228203).

The current OASIS3-MCT version was significantly refactored with respect to OASIS3.3. OASIS3-MCT is now interfaced with the Model Coupling Toolkit¹ (**author?**) (Larson et al 2005) (**author?**) (Jacob et al 2005) developed by the Argonne National Laboratory in the USA. MCT implements fully parallel regridding (as a parallel matrix vector multiplication) and parallel distributed exchanges of the coupling fields, based on pre-computed regridding weights and addresses. Its design philosophy, based on flexibility and minimal invasiveness, is close to the OASIS approach. MCT has proven parallel performance and is, most notably, the underlying coupling software used in National Center for Atmospheric Research Community Earth System Model 1 (NCAR CESM1).

OASIS3-MCT is a portable set of Fortran 77, Fortran 90 and C routines. Low-intrusiveness, portability and flexibility are OASIS3-MCT key design concepts. At run-time, there is no longer a separate coupler executable: OASIS3-MCT acts as a coupling library, which main function is to interpolate and exchange the coupling fields between the component models of a coupled system. OASIS3-MCT supports coupling of general two dimensional fields. Unstructured grids are also supported using a one dimension degeneration of the two dimensional structures. Thanks to MCT, all transformations, including regridding, are executed in parallel on the set of source or target component processes and all couplings are now executed in parallel directly between the components via Message Passing Interface (MPI). OASIS3-MCT also supports parallel file I/O using netcdf.

In spite of the significant changes in underlying implementation, usage of OASIS3-MCT in the codes has largely remained unchanged with respect to OASIS3.3. To communicate with another model, or to perform I/O actions, a component model needs to include few specific calls to the OASIS3-MCT coupling library, which Application Programmig Interface used in component models is unchanged. The use statement has been updated and now requires a single “use mod_prism” or “use mod_oasis” statement instead of the various use statements required in prior OASIS3 versions. The *namcouple* configuration file is also largely unchanged relative to OASIS3, although several options are either not used or not supported. There is a new transformation in *namcouple* i.e. MAPPING which allows a user to specify a mapping file generated externally. Some features like vector mapping and second order mapping have been delayed in implementation while other new features like parallel mapping have been added. And currently, only

¹MCT, see www.mcs.anl.gov/research/projects/mct/

MPI1 job launching is supported.

First tests done with up to 8000 cores on the Bullx Curie machine at the TGCC are very encouraging and it is therefore very likely that OASIS3-MCT will provide an efficient and easy-to-use coupling solution.

1.1 Step-by-step use of OASIS3-MCT

To use OASIS3-MCT for coupling models (and/or perform I/O actions), one has to follow these steps:

1. Obtain OASIS3-MCT source code (see chapter 1.2).
2. Identify the coupling or I/O fields and adapt the component models to allow their exchange with the OASIS3-MCT coupling library based on MPI1 message passing. The OASIS3-MCT coupling library uses NetCDF and therefore can be used to perform I/O actions from/to disk files. For more detail on how to interface a model with OASIS3-MCT, see chapter 2.

The tutorial coupled model gives a practical example of a coupled model; the sources are given in directories `examples/tutorial` ; more detail on the `tutorial` and how to compile and run it can be found in chapter 6.

3. Define all coupling and I/O parameters and the transformations required to adapt each coupling field from its source model grid to its target model grid; on this basis, prepare OASIS3-MCT configuring file *namcouple* (see chapter 3).

OASIS3-MCT supports different interpolation algorithms as is described in chapter 4. Regridding files can be compute online using the SCRIP options or offline and read using the MAPPING transformation.

4. Generate required auxiliary data files (see chapter 5) XXXXX.
5. Compile OASIS3-MCT, the component models and start the coupled experiment. Chapter 6 describes how to compile and run OASIS3-MCT and the `tutorial` toy coupled model.

If you need extra help, do not hesitate to contact us (see contact details on the back of the cover page).

1.2 OASIS3-MCT sources

OASIS3-MCT and `tutorial` toy coupled model sources are available from CERFACS SVN server. To obtain more detail on how to download the sources, please fill in the registration form at <https://verc.enes.org/oasis/download/oasis-registration-form> .

OASIS3-MCT directory structure is the following one:

- `oasis3-mct/lib/psmile` OASIS3-MCT coupling library
- `/scrip` SCRIP interpolation library
- `/mct` Model Coupling Toolkit Coupling Software
- `oasis3-mct/doc` OASIS3-MCT User Guide
- `oasis3-mct/util/make_dir` Utilities to compile OASIS3-MCT
- `oasis3-mct/examples/tutorial` Environment to run the tutorial toy model

1.3 Licenses and Copyrights

1.3.1 OASIS3-MCT license and copyright statement

Copyright 2012 Centre Européen de Recherche et Formation Avancée en Calcul Scientifique (CERFACS).

This software and ancillary information called OASIS3-MCT is free software. CERFACS has rights to use, reproduce, and distribute OASIS3-MCT. The public may copy, distribute, use, prepare derivative works and publicly display OASIS3-MCT under the terms of the Lesser GNU General Public License (LGPL) as published by the Free Software Foundation, provided that this notice and any statement of authorship are reproduced on all copies. If OASIS3-MCT is modified to produce derivative works, such modified software should be clearly marked, so as not to confuse it with the OASIS3-MCT version available from CERFACS.

The developers of the OASIS3-MCT software are researchers attempting to build a modular and user-friendly coupler accessible to the climate modelling community. Although we use the tool ourselves and have made every effort to ensure its accuracy, we can not make any guarantees. We provide the software to you for free. In return, you—the user—assume full responsibility for use of the software. The OASIS3-MCT software comes without any warranties (implied or expressed) and is not guaranteed to work for you or on your computer. Specifically, CERFACS and the various individuals involved in development and maintenance of the OASIS3-MCT software are not responsible for any damage that may result from correct or incorrect use of this software.

1.3.2 MCT copyright statement

Modeling Coupling Toolkit (MCT) Software

Copyright 2011, UChicago Argonne, LLC as Operator of Argonne National Laboratory. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
3. The end-user documentation included with the redistribution, if any, must include the following acknowledgment: "This product includes software developed by the UChicago Argonne, LLC, as Operator of Argonne National Laboratory." Alternately, this acknowledgment may appear in the software itself, if and wherever such third-party acknowledgments normally appear.

This software was authored by:

- Argonne National Laboratory Climate Modeling Group, Mathematics and Computer Science Division, Argonne National Laboratory, Argonne IL 60439
 - Robert Jacob, tel: (630) 252-2983, E-mail: jacob@mcs.anl.gov
 - Jay Larson, E-mail: larson@mcs.anl.gov
 - Everest Ong
 - Ray Loy
4. **WARRANTY DISCLAIMER. THE SOFTWARE IS SUPPLIED "AS IS" WITHOUT WARRANTY OF ANY KIND. THE COPYRIGHT HOLDER, THE UNITED STATES, THE UNITED STATES DEPARTMENT OF ENERGY, AND THEIR EMPLOYEES: (1) DISCLAIM ANY WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, TITLE OR NON-INFRINGEMENT, (2) DO NOT ASSUME ANY LEGAL LIABILITY OR RESPONSIBILITY FOR THE ACCURACY, COMPLETENESS, OR USEFULNESS OF THE SOFTWARE, (3) DO NOT REPRESENT THAT USE OF THE SOFTWARE WOULD NOT INFRINGE PRIVATELY OWNED RIGHTS, (4) DO NOT WARRANT THAT THE SOFTWARE WILL FUNCTION UNINTERRUPTED, THAT IT IS ERROR-FREE OR THAT ANY ERRORS WILL BE CORRECTED.**

5. LIMITATION OF LIABILITY. IN NO EVENT WILL THE COPYRIGHT HOLDER, THE UNITED STATES, THE UNITED STATES DEPARTMENT OF ENERGY, OR THEIR EMPLOYEES: BE LIABLE FOR ANY INDIRECT, INCIDENTAL, CONSEQUENTIAL, SPECIAL OR PUNITIVE DAMAGES OF ANY KIND OR NATURE, INCLUDING BUT NOT LIMITED TO LOSS OF PROFITS OR LOSS OF DATA, FOR ANY REASON WHATSOEVER, WHETHER SUCH LIABILITY IS ASSERTED ON THE BASIS OF CONTRACT, TORT (INCLUDING NEGLIGENCE OR STRICT LIABILITY), OR OTHERWISE, EVEN IF ANY OF SAID PARTIES HAS BEEN WARNED OF THE POSSIBILITY OF SUCH LOSS OR DAMAGES.

1.3.3 The SCRIP 1.4 license copyright statement

The SCRIP 1.4 copyright statement reads as follows:

“Copyright 1997, 1998 the Regents of the University of California. This software and ancillary information (herein called SOFTWARE) called SCRIP is made available under the terms described here. The SOFTWARE has been approved for release with associated LA-CC Number 98-45. Unless otherwise indicated, this SOFTWARE has been authored by an employee or employees of the University of California, operator of Los Alamos National Laboratory under Contract No. W-7405-ENG-36 with the United States Department of Energy. The United States Government has rights to use, reproduce, and distribute this SOFTWARE. The public may copy, distribute, prepare derivative works and publicly display this SOFTWARE without charge, provided that this Notice and any statement of authorship are reproduced on all copies. Neither the Government nor the University makes any warranty, express or implied, or assumes any liability or responsibility for the use of this SOFTWARE. If SOFTWARE is modified to produce derivative works, such modified SOFTWARE should be clearly marked, so as not to confuse it with the version available from Los Alamos National Laboratory.”

Chapter 2

Interfacing a component model with the OASIS3-MCT library

At run-time, the OASIS3-MCT coupling layer supports coupling data between two components as well as interpolation and transformation of coupling fields. Different communication techniques have been historically developed in OASIS. With OASIS3-MCT, communication is performed by MCT based on Message Passing Interface 1 protocol (the keyword `$CHANNEL` in the configuration file *namcouple* has to be `MP11`, see chapter 3).

For a practical test case using the OASIS3-MCT library, see the sources in `examples/tutorial` and more details in chapter 6.

To communicate with another component model or to perform I/O actions, a component model needs to be interfaced with the OASIS3-MCT library, which sources can be found in `oasis3-mct/lib/psmile` directory. The OASIS3-MCT library supports:

- parallel communication between parallel component models,
- an ability to couple a component on a subset of its processes only,
- automatic sending and receiving actions at appropriate times following user's choice indicated in the *namcouple*,
- time integration or accumulation of the coupling fields,
- some transformations such as mapping (interpolation) between grids,
- I/O actions from/to files.

To adapt a component model to OASIS3-MCT, specific calls of the following classes have to be implemented in the code:

1. Initialisation (section 2.2)
2. Grid data file definition (section 2.3)
3. Partition definition (section 2.4)
4. coupling-I/O field declaration (section 2.5)
5. End of definition phase (section 2.6)
6. coupling-I/O field sending and receiving (section 2.7)
7. Termination (section 2.8)
8. Optional auxiliary routines (section 2.9)

Finally, in section 2.10, different coupling algorithms are illustrated and details on how to reproduce them with OASIS3-MCT are provided. More information on the `LAG` and `SEQ` indices are also given in that section.

2.1 Use of OASIS3-MCT library

To use OASIS3-MCT library, a user needs to add in his code:

- `USE mod_oasis`
 `** OR **`
- `USE mod_prism`

Both use statements are valid and use of just one or the other is recommended in a particular component model. A single use statement now provides all the methods that required multiple use statements in previous OASIS3 versions. The methods, datatypes, and capabilities are identical for both the `mod_prism` or `mod_oasis` interfaces. The only difference is the name of the interface. The interface in module `mod_prism` is provided for backwards compatability with prior versions of OASIS3. Use of module `mod_oasis` is now recommended provides access to a set of updated routine names that will continue to evolve in the future, always ensuring backward compatibility. In the following sections, both the `mod_prism` and `mod_oasis` interface names is defined and a single description of the interface arguments is provided.

2.2 Initialisation

2.2.1 Coupling initialisation

- `CALL oasis_init_comp (compid, model_name, ierror)`
- `CALL prism_init_comp_proto (compid, model_name, ierror)`
 - `compid` [INTEGER; OUT]: component model ID
 - `model_name` [CHARACTER*6; IN]: component model name (as in *namcouple*)
 - `ierror` [INTEGER; OUT]: returned error code.

This routine must called by all component processes to initialise the coupling.¹

2.2.2 Communicator for internal parallelisation

- `CALL oasis_get_localcomm (local_comm, ierror)`
- `CALL prism_get_localcomm_proto (local_comm, ierror)`
 - `local_comm` [INTEGER; OUT]: value of local communicator
 - `ierror` [INTEGER; OUT]: returned error code.

If needed, this routine may be called by the component processes to get the value of a local communicator to be used by the component for its internal parallelisation.

This may be needed as all component models started in a pseudo-MPMD mode with MPI1 share automatically the same `MPI_COMM_WORLD` communicator. Another communicator has to be used for the internal parallelisation of each component. OASIS3-MCT creates this local communicator based on the name given to `oasis_init_comp/prism_init_comp_proto` routine; its value is returned as the first argument of the routine, `local_comm`.

¹The model may call `MPI_Init` explicitly, but if so, has to call it before calling `prism_init_comp_proto`; in this case, the model also has to call `MPI_Finalize` explicitly, but only after calling `prism_terminate_proto`.

2.2.3 Coupling through a subset of the component model processes

If only a subset of the component processes participate in the coupling (e.g. a component is setup to run on 80 processes but 16 of those processes are associated with a distinct task, like I/O), a communicator gathering only these processes must be defined, with either `oasis/prism.create_couplcomm` or `oasis/prism.set_couplcomm`.

If such communicator does not exist yet in the code, the component processes should use, to create it:

- CALL `oasis.create_couplcomm(icpl, local_comm, coupl_comm, kinfo)`
- CALL `prism.create_couplcomm(icpl, local_comm, coupl_comm, kinfo)`
 - `icpl` [INTEGER; IN]: coupling process flag
 - `local_comm` [INTEGER; IN]: MPI communicator with all processes of the component
 - `coupl_comm` [INTEGER; OUT]: returned MPI communicator gathering only component processes participating in the coupling
 - `kinfo` [INTEGER; OUT; OPTIONAL]: returned error code

This routine creates a coupling communicator for a subset of processes. It must be called by all component processes with `icpl=1` for processes participating in the coupling and with `icpl=MPI_UNDEFINED` for the others. Argument `local_comm` is the MPI communicator associated with all processes of the component. The new coupling communicator is returned in `coupl_comm` argument and the internal coupling communicator is also set to that value.

If this communicator already exist in the code, the component should use, to provide it to OASIS3-MCT:

- CALL `oasis.set_couplcomm(coupl_comm, kinfo)`
- CALL `prism.set_couplcomm(coupl_comm, kinfo)`
 - `coupl_comm` [INTEGER; IN]: MPI communicator gathering only component processes participating in the coupling
 - `kinfo` [INTEGER; OUT; OPTIONAL]: returned error code

This routine allows to provide a local coupling communicator to OASIS3-MCT, given that it already exists in the code. The value of `coupl_comm` must be the value of this local coupling communicator for the processes participating to the coupling and it must be `MPI_COMM_NULL` for processes not involved in the coupling.

These routines should be called after the `oasis_init_comp/prism_init_comp_proto` but before the grid, partition, or coupling field declaration. All OASIS3-MCT interface routines, besides the grid definition (see section 2.3) and the “puts” and “gets” per se (see section 2.7), are collective and must be called by all processes. In particular, the coupling fields sent or received by the component must be declared by all component processes with `oasis_def_var/prism_def_var_proto` (see section 2.5) and the field partition must be described across all coupling processes with `oasis_def_partition/prism_def_partition_proto` (but with `ig_parallel(:)=0` for the processes not involved in the coupling, see section 2.4.)

Here is a coding sample of how to use these routines:

```
CALL oasis_init_comp (comp_id, comp_name, ierror )
CALL oasis_get_localcomm ( localComm, ierror )

!--- create communicator gathering coupling processes (every other)
CALL MPI_Comm_Rank ( localComm, mype, ierror )
couplingpe = .false.
if (mod(mype,2) == 0) couplingpe = .true.
icpl = MPI_UNDEFINED
if (couplingpe) icpl = 1
```

```

CALL MPI_COMM_Split(localComm,icpl,1,couplComm,ierror)
!
!--- provide this communicator to OASIS3-MCT
CALL oasis_set_couplcomm(couplComm, ierror)

! The call to MPI_COMM_Split and oasis_set_couplcomm could be replaced by
! CALL oasis_create_couplcomm(icpl,localComm,couplComm,ierror)

CALL oasis_def_partition ( ... )
CALL oasis_def_var ( ... )
CALL oasis_enddef ( ... )

!--- do loop
! ...
if (couplingpe) CALL oasis_put( ... )
! ...
if (couplingpe) CALL oasis_get( ... )
! ...
!--- enddo

CALL oasis_terminate ( ... )

```

2.3 Grid data file definition

With OASIS3-MCT, the grid data files *grids.nc*, *masks.nc* and *areas.nc* are required only for certain operations (see also section 5.2), i.e. *grids.nc*, and *masks.nc* for SCRIPR (see section 4.3) and *masks.nc* and *areas.nc* for CONSERV (see section 4.4). These grid data files can be created by the user before the run or can be written directly at run time by the **master process of each component model** with the following routines. These routines can be used by the component models to add grid fields to the grid files but grid fields are **never** overwritten in the grid files. These routines have to be called just after `prism_init_comp`.

- CALL `oasis_start_grids_writing (flag)` or
- CALL `prism_start_grids_writing (flag)`

– flag [INTEGER; OUT]: not used

Obsolete in OASIS3-MCT; exists however for upward compatibility.

- CALL `oasis_write_grid (cgrid, nx, ny, lon, lat)` or
- CALL `prism_write_grid (cgrid, nx, ny, lon, lat)`

– cgrid [CHARACTER*4; IN]: grid name prefix (see 3.3)

– nx [INTEGER; IN] : first grid dimension (x)

– ny [INTEGER; IN] : second grid dimension (y)

– lon [REAL, DIMENSION(nx,ny); IN] : array of longitudes (degrees East)

– lat [REAL, DIMENSION(nx,ny); IN] : array of latitudes (degrees North)

Writes the model grid longitudes and latitudes. Longitudes must be given in degrees East in the interval -360.0 to 720.0. Latitudes must be given in degrees North in the interval -90.0 to 90.0. Note that if some grid points overlap, it is recommended to define those points with the same number (e.g. 90.0 for both, not 450.0 for one and 90.0 for the other) to ensure automatic detection of overlap by

OASIS3-MCT (which is essential to have a correct conservative remapping SCRIPR/CONSERV, see section 4.3).

- CALL oasis_write_corner (cgrid, nx, ny, nc, clon, clat) or
- CALL prism_write_corner (cgrid, nx, ny, nc, clon, clat)
 - cgrid [CHARACTER*4; IN]: grid name prefix
 - nx [INTEGER; IN] : first grid dimension (x)
 - ny [INTEGER; IN] : second grid dimension (y)
 - nc [INTEGER; IN] : number of corners per grid cell (always 4 in the version)
 - lon [REAL, DIMENSION (nx,ny,nc) ; IN] : array of corner longitudes (in degrees East)
 - lat [REAL, DIMENSION (nx,ny,nc) ; IN] : array of corner latitudes (in degrees North)

Writes the grid cell corner longitudes and latitudes (counterclockwise sense). Longitudes must be given in degrees East in the interval -360.0 to 720.0. Latitudes must be given in degrees North in the interval -90.0 to 90.0. Note also that cells larger than 180.0 degrees in longitude are not supported. Writing of corners is optional as corner information is needed only for SCRIPR/CONSERV (see section 4.3). If called, needs to be called after oasis/prism_write_grid.

- CALL oasis_write_angle (cgrid, nx, ny, angle) or
- CALL prism_write_angle (cgrid, nx, ny, angle)
 - cgrid [CHARACTER*4; IN]: grid name prefix
 - nx [INTEGER; IN] : first grid dimension (x)
 - ny [INTEGER; IN] : second grid dimension (y)
 - angle [REAL, DIMENSION (nx,ny) ; IN] : array of angles

Obsolete in OASIS3-MCT as vector interpolation is not supported. See SCRIPR/CONSERV in section 4.3.

- CALL oasis_write_mask (cgrid, nx, ny, mask) or
- CALL prism_write_mask (cgrid, nx, ny, mask)
 - cgrid [CHARACTER*4; IN]: grid name prefix
 - nx [INTEGER; IN] : first grid dimension (x)
 - ny [INTEGER; IN] : second grid dimension (y)
 - mask [INTEGER, DIMENSION (nx,ny) ; IN] : mask array (be careful about the OASIS historical convention (!): 0 = not masked, 1 = masked)

Writes the model grid mask.

- CALL oasis_write_area (cgrid, nx, ny, area) or
- CALL prism_write_area (cgrid, nx, ny, area)
 - cgrid [CHARACTER*4; IN]: grid name prefix
 - nx [INTEGER; IN] : first grid dimension (x)
 - ny [INTEGER; IN] : second grid dimension (y)
 - area [REAL, DIMENSION (nx,ny) ; IN] : array of grid cell areas

Writes of the model grid cell areas. Needed only for CONSERV operation (see section 4.4).

- CALL prism_terminate_grids_writing()
- CALL oasis_terminate_grids_writing()

Terminates grids writing (required if some of the above grid writing routines are called).

2.4 Partition definition

The coupling fields sent or received by a component model are usually scattered among the different component processes. With OASIS3-MCT, all processes exchanging coupling data have to define their local partition in the global index space. If only a subset of the processes actually exchange coupling data, the processes not implied in the coupling have to call this routine to describe a null partition (i.e. with `ig_paral(:)=0`).

- CALL `oasis.def.partition (il_part_id, ig_paral, ierror)` or
- CALL `prism.def.partition_proto (il_part_id, ig_paral, ierror)`
 - `il_part_id` [INTEGER; OUT]: partition ID
 - `ig_paral` [INTEGER, DIMENSION(:), IN]: vector of integers describing the local partition in the global index space; has a different expression depending on the type of the partition; in OASIS3-MCT, 4 types of partition are supported: Serial (no partition), Apple, Box, and Orange (see below).
 - `ierror` [INTEGER; OUT]: returned error code.

2.4.1 Serial (no partition)

This is the choice for a monoprocess model. In this case, we have `ig_paral(1:3):`

- `ig_paral(1) = 0` (indicates a Serial “partition”)
- `ig_paral(2) = 0`
- `ig_paral(3) = the total grid size.`

2.4.2 Apple partition

Each partition is a segment of the global domain, described by its global offset and its local size. In this case, we have `ig_paral(1:3):`

- `ig_paral(1) = 1` (indicates an Apple partition)
- `ig_paral(2) = the segment global offset`
- `ig_paral(3) = the segment local size`

Figure 2.1 illustrates an Apple partition over 3 processes.

2.4.3 Box partition

Each partition is a rectangular region of the global domain, described by the global offset of its upper left corner, and its local extents in the X and Y dimensions. The global extent in the X dimension must also be given. In this case, we have `ig_paral(1:5):`

- `ig_paral(1) = 2` (indicates a Box partition)
- `ig_paral(2) = the upper left corner global offset`
- `ig_paral(3) = the local extent in x`
- `ig_paral(4) = the local extent in y`
- `ig_paral(5) = the global extent in x.`

Figure 2.2 illustrates a Box partition over 3 processes.

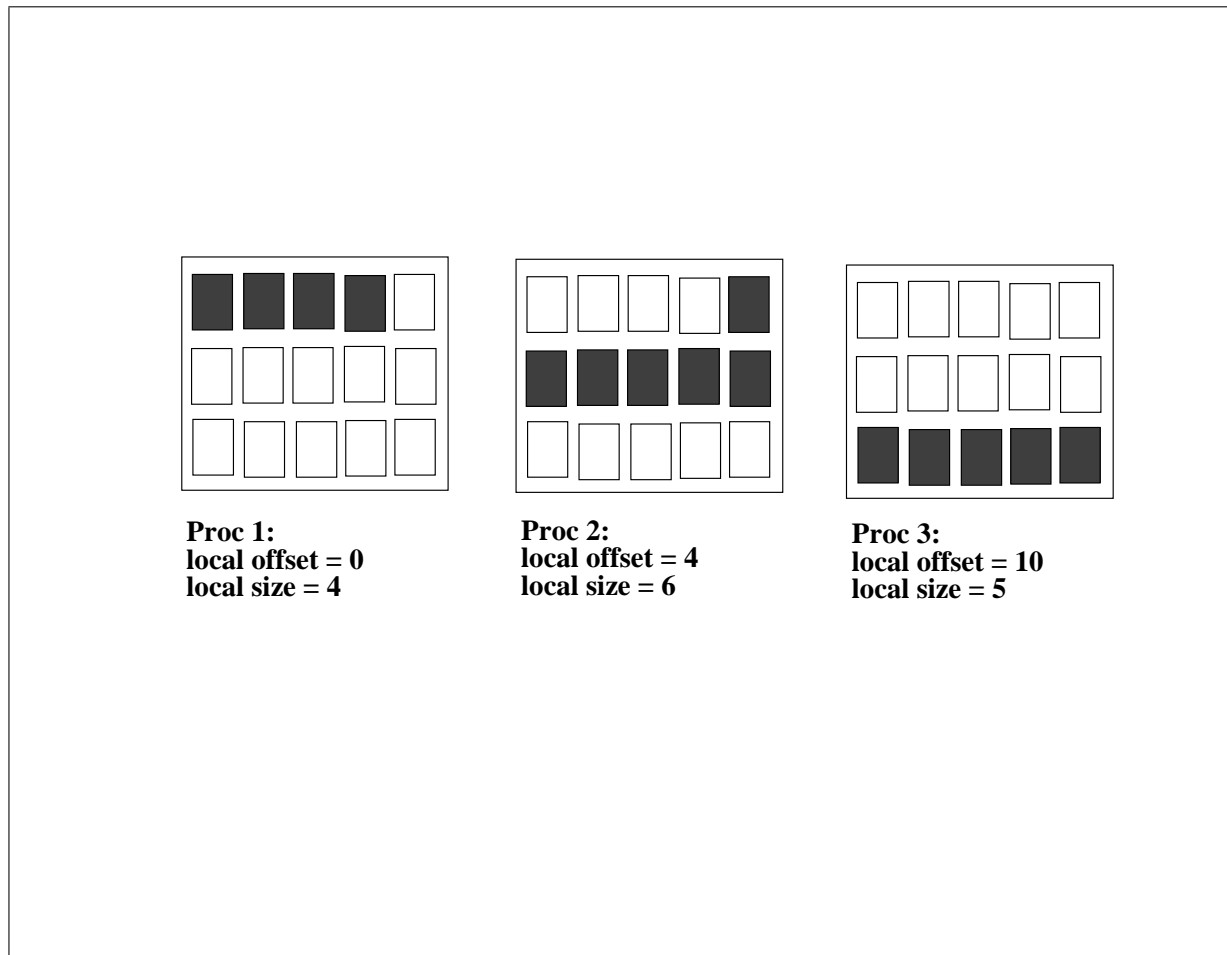


Figure 2.1: Apple partition. It is assumed here that the index start at 0 in the upper left corner.




Figure 2.2: Box partition. It is assumed here that the index start at 0 in the upper left corner.




Figure 2.3: Orange partition for one process. It is assumed here that the index start at 0 in the upper left corner.

2.4.4 Orange partition

Each partition is an ensemble of segments of the global domain. Each segment is described by its global offset and its local extent. In this case, we have `ig_paral(1:N)` where $N = 2 + 2 \times \text{number of segments}$

- `ig_paral(1) = 3` (indicates a Orange partition)
- `ig_paral(2)` = the total number of segments for the partition (limited to 200 presently, see note for `ig_paral(4)` for Box partition above)
- `ig_paral(3)` = the first segment global offset
- `ig_paral(4)` = the first segment local extent
- `ig_paral(5)` = the second segment global offset
- `ig_paral(6)` = the second segment local extent
- ...
- `ig_paral(N-1)` = the last segment global offset
- `ig_paral(N)` = the last segment local extent

Figure 2.3 illustrates an Orange partition with 3 segments for one process. The other process partitions are not illustrated.

2.5 Coupling field declaration

All component processes declares the fields sent or received by the component during the simulation. This is true even for processes not implied in the coupling.

- CALL oasis_def_var (var_id, name, il_part_id, var_nodims, kinout, var_actual_shape, var_type, ierror) or
- CALL prism_def_var_proto(var_id, name, il_part_id, var_nodims, kinout, var_actual_shape, var_type, ierror)
 - var_id [INTEGER; OUT]: coupling field ID
 - name [CHARACTER*8; IN]: field symbolic name (as in the *namcouple*)
 - il_part_id [INTEGER; IN]: partition ID (see section 2.4)
 - var_nodims [INTEGER, DIMENSION(2); IN]: var_nodims(1) is the rank of field array (1 or 2); var_nodims(2) is the number of bundles (always 1 for OASIS3-MCT).
 - kinout [INTEGER; IN]: OASIS_In or PRISM_In (i.e. = 15) for fields received by the model; OASIS_Out, PRISM_Out (i.e. = 14) for fields sent by the model ².
 - var_actual_shape [INTEGER, DIMENSION(2*var_nodims(1)); IN]: vector of integers giving the minimum and maximum index for each dimension of the coupling field array; for OASIS3, the minimum index has to be 1 and the maximum index has to be the extent of the dimension.
 - var_type [INTEGER; IN]: type of coupling field array; put OASIS_Real or PRISM_Real (i.e. = 4) for single or double precision real arrays. All coupling data is treated as double precision in the coupling layer, but conversion to or from single precision data is supported in the interface.
 - ierror [INTEGER; OUT]: returned error code.

2.6 End of definition phase

All component processes close the definition phase.

- CALL oasis_enddef (ierror)
- CALL prism_enddef_proto(ierror)
 - ierror [INTEGER; OUT]: returned error code.

2.7 Sending and receiving actions

2.7.1 Sending a coupling (or I/O) field

In the model time stepping loop, each process sends its part of the coupling (or I/O) field.

- CALL oasis_put (var_id, date, field_array, info)
- CALL prism_put_proto(var_id, date, field_array, info)
 - var_id [INTEGER; IN]: field ID (from corresponding oasis_def_var / prism_def_var_proto, see section 2.5)
 - date [INTEGER; IN]: number of seconds in the run at the time of the call (by convention at the beginning of the timestep)
 - field_array [REAL, IN]: coupling (or I/O) field array
 - info [INTEGER; OUT]: returned info code:
 - * OASIS_Sent(=4) if the field was sent to another model
 - * OASIS_LocTrans(=5) if the field was only used in a time transformation (not sent, not output)

²Parameters OASIS_In, PRISM_In, OASIS_Out, PRISM_Out are defined in oasis3-mct/lib/psmile/src/mod_oasis_parameters.F90

- * OASIS_ToRest (=6) if the field was written to a restart file only
- * OASIS_Output (=7) if the field was written to an output file only
- * OASIS_SentOut (=8) if the field was both written to an output file and sent to another model (directly or via OASIS3 main process)
- * OASIS_ToRestOut (=9) if the field was written both to a restart file and to an output file.
- * OASIS_Ok (=0) otherwise and no error occurred.

This routine may be called by the model at each timestep. The sending is actually performed only if the time obtained by adding the field lag (`LAG` in the *namcouple*) to the argument `date` corresponds to a time at which it should be activated, given the coupling or I/O period indicated by the user in the *namcouple* (see section 3). A field will not be sent at all if its coupling (or I/O) period indicated in the *namcouple* is greater than the total run time.

An exchange at the beginning of the run at time = 0 is expected; if a coupling field has a positive lag, the coupling field that matches the `oasis_get / prism_get_proto` at time=0 will come from a coupling restart file written by the last `oasis_put / prism_put_proto` of the previous run (see section 2.10). For a coupling field with a positive lag, the coupling restart file (see section 5.3) is automatically written by the last `prism_put_proto` call of the run, if the time = `date+LAG` corresponds to a coupling or I/O period.

If a local time transformation is indicated for the field by the user in the *namcouple* (INSTANT, AVERAGE, ACCUMUL, T_MIN or T_MAX, see section 4), it is automatically performed and the resulting field is finally sent at the coupling or I/O frequency. For non-instant transformations, partially transformed fields will be written to the restart file at the end of the run for use on the next model startup; this is a bug fix new in OASIS3-MCT.

2.7.2 Receiving a coupling (or I/O) field

In the model time stepping loop, each process receives its part of the coupling field.

- CALL `oasis_get (var_id, date, field_array, ierror)`
- CALL `prism_get_proto(var_id, date, field_array, ierror)`
 - `var_id` [INTEGER; IN]: field ID (from corresponding `prism_def.var_proto`)
 - `date` [INTEGER; IN]: number of seconds in the run at the time of the call
 - `field_array` [REAL, OUT]: I/O or coupling field array
 - `info` [INTEGER; OUT]: returned info code:
 - * OASIS_Recvd(=3) if the field was received from another model
 - * OASIS_FromRest (=10) if the field was read from a restart file only
 - * OASIS_Input (=11) if the field was read from an input file only
 - * OASIS_RecvOut (=12) if the field was both received from another model and written to an output file
 - * OASIS_FromRestOut (=13) if the field was both read from a restart file and written to an output file
 - * OASIS_Ok (=0) otherwise and no error occurred.

This routine may be called by the model at each timestep. The `date` argument is automatically analysed and the receiving action is actually performed only if `date` corresponds to a time for which it should be activated, given the period indicated by the user in the *namcouple*. An exchange at the beginning of the run at time=0 is expected. A field will not be received at all if its coupling or I/O period indicated in the *namcouple* is greater than the total run time.

2.8 Termination

- CALL `oasis.terminate (ierror)`
- CALL `prism.terminate_proto(ierror)`
 - `ierror [INTEGER; OUT]`: returned error code.

All processes of the component model must terminate the coupling by calling this routine³ (normal termination).

2.9 Auxiliary routines

Not all auxiliary routines that were in OASIS3.3 are currently available.

- CALL `oasis.abort (compid, routine_name, abort_message)`
- CALL `prism.abort_proto(compid, routine_name, abort_message)`
 - `compid [INTEGER; IN]`: component model ID (from `oasis.init_comp` or `prism.init_comp_proto`)
 - `routine_name[CHARACTER*; IN]`: name of calling routine
 - `abort_message[CHARACTER*; IN]`: message to be written out.

If a process needs to abort voluntarily, it should do so by calling `oasis.abort/prism.abort_proto`. This will ensure a proper termination of all processes in the coupled model communicator. This routine writes the name of the calling model, the name of the calling routine, and the message to the process debug file (see `$NLOGPRT` in section 3.2). This routine cannot be called before `prism.init_comp_proto`.

- CALL `oasis.get_debug(debug_value)`
- CALL `prism.get_debug(debug_value)`
 - `debug_value [INTEGER; OUT]`: debug value

This routine may be called at any time to retrieve the current OASIS3-MCT internal debug level (see `$NLOGPRT` in section 3.2). This is useful if the user wants to return the original debug value after changing it or if a user wants to key off the oasis debug level for model debug diagnostics.

- CALL `oasis.set_debug(debug_value)`
- CALL `prism.set_debug(debug_value)`
 - `debug_value [INTEGER; IN]`: debug value

This routine may be called at any time to change the debug level in oasis. This method allows users to vary the debug level at different points in the model integration.

- CALL `oasis.get_intercomm(new_comm, cdnam, kinfo)`
- CALL `prism.get_intercomm(new_comm, cdnam, kinfo)`
 - `new_comm [INTEGER; OUT]`: mpi intercomm communicator
 - `cdnam [CHARACTER*; IN]`: other model name
 - `kinfo [INTEGER; OUT; OPTIONAL]`: returned error code

This routine sets up an MPI intercomm communicator between the root processors of two components, the local component and the component associated with `cdnam`. This must be called by both components at the same time otherwise a deadlock will occur. In addition, this call is collective across the tasks of the two components but other components are not involved.

³If the process called `MPI.Init` (before calling `oasis.init_comp` or `prism.init_comp_proto`), it must also call `MPI.Finalize` explicitly, but only after calling `oasis.terminate_proto` or `prism.terminate_proto`.

- CALL `oasis_get_intracomm(new_comm, cdnam, kinfo)`
- CALL `prism_get_intracomm(new_comm, cdnam, kinfo)`
 - `new_comm` [INTEGER; OUT]: mpi intracomm communicator
 - `cdnam` [CHARACTER*; IN]: other model name
 - `kinfo` [INTEGER; OUT; OPTIONAL]: returned error code

This routine sets up an MPI intracomm communicator between the root processors of two components, the local component and the component associated with `cdnam`. This must be called by both components at the same time otherwise a deadlock will occur. In addition, this call is collective across the tasks of the two components but other components are not involved.

2.10 Coupling algorithms - SEQ and LAG concepts

Using the OASIS3-MCT coupling library, the user has full flexibility to reproduce different coupling algorithms. In the components, the sending and receiving routines, respectively `oasis_put/prism_put_proto` and `oasis_get/prism_get_proto`, can be called at each model timestep, with the appropriate `date` argument giving the actual time (at the beginning of the timestep), expressed in “number of seconds since the start of the run” (see section 2.7.1). This `date` argument is automatically analysed by the coupling library and depending on the coupling period and the lag (LAG) chosen by the user for each coupling field in the configuration file *namcouple*, different coupling algorithms can be reproduced without modifying anything in the component model codes themselves.

With OASIS3-MCT, the SEQ index is no longer needed in the *namcouple* input to specify the sequencing order of different coupling fields. The sequence (SEQ) index in the *namcouple* file now provides the coupling layer with an ability to detect a deadlock before it happens and exit.

The lag concept and indices are explained in more detail below and some examples are provided.

2.10.1 The lag concept

The lag (LAG) value tells the coupler to modify the time at which that data is sent (put) by the amount of the lag. The lag must be expressed in “number of seconds” and can be positive or negative but should never be larger (in absolute magnitude) than the coupling period of any field due to problems with restart-ability and dead-locking. When a component model calls a `oasis_put/prism_put_proto`, the value of the lag is automatically added to the value of the `date` argument and the “put” is actually performed when the sum `date+lag` is a coupling time; in the target component, this “put” will match a `oasis_get/prism_get_proto` for which the `date` argument is the same coupling time. The lag only shifts the time data is sent. It cannot be used to shift the time data is received yet.

When the lag is positive, a restart file must be available to initiate the coupling and in those cases, the restart file is then updated at the end of the run. A positive lag acts like a send occurred before the model started. In fact, for a field with positive lag, the source component model automatically reads the field in the restart file during the coupling initialization phase (below the `oasis_enddef/prism_enddef_proto`) and send the data to match the `oasis_get/prism_get_proto` performed at `time=0` in the target component model. The final coupling data on the source side will then be automatically written to the restart file for use in the next run.

When there is a lag, the first and last instance of the source field in the debug netCDF file (EXPOUT fields, see section 3.3) always correspond respectively to the field read from and written to the restart file.

1. LAG concept first example

A first coupling algorithm, exploiting the LAG concept, is illustrated on figure 2.4.

On the 4 figures in this section, short black arrows correspond to `oasis_put/prism_put_proto` or `oasis_get/prism_get_proto` called in the component model that do not lead to any “put”




Figure 2.4: LAG concept first example

or receiving action; long black arrows correspond to `oasis_put/prism_put_proto` or `oasis_get/prism_get_proto` called in the component models that do actually lead to a “put” or “get” action; long red arrows correspond to `oasis_put/prism_put_proto` or `oasis_get/prism_get_proto` called in the component models that lead to a reading or writing of the coupling field from or to a coupling restart file.

During a coupling timestep, model A receives F_2 and then sends F_1 ; its timestep length is 4. During a coupling timestep, model B receives F_1 and then sends F_2 ; its timestep length is 6. F_1 and F_2 coupling periods are respectively 12 and 24. If F_1/F_2 “put” action by model A/B was used at a coupling timestep to match the model B/A “get” action, a deadlock would occur as both models would be initially waiting on a “get” action. To prevent this, F_1 and F_2 produced at the timestep before have to be used to match respectively the model B and model A “get” actions.

This implies that a lag of respectively 4 and 6 seconds must be defined for F_1 and F_2 . For F_1 , the `oasis_put/prism_put_proto` performed at time 8 and 20 by model A will then lead to “put” actions (as $8 + 4 = 12$ and $20 + 4 = 24$ which are coupling periods) that match the “get” actions performed at times 12 and 24 below the `oasis_get/prism_get_proto` called by model B. For F_2 , the `oasis_put/prism_put_proto` performed at time 18 by model B then leads to a “put” action (as $18 + 6 = 24$ which is a coupling period) that matches the “get” action performed at time 24 below the `oasis_get/prism_get_proto` called by model A.

At the beginning of the run, as their LAG index is greater than 0, the first `oasis_get/prism_get_proto` of F_1 and F_2 will automatically be fulfilled with fields read from their respective coupling restart files. The user therefore has to create those coupling restart files before the first run in the experiment. At the end of the run, F_1 having a lag greater than 0, is automatically written to its coupling

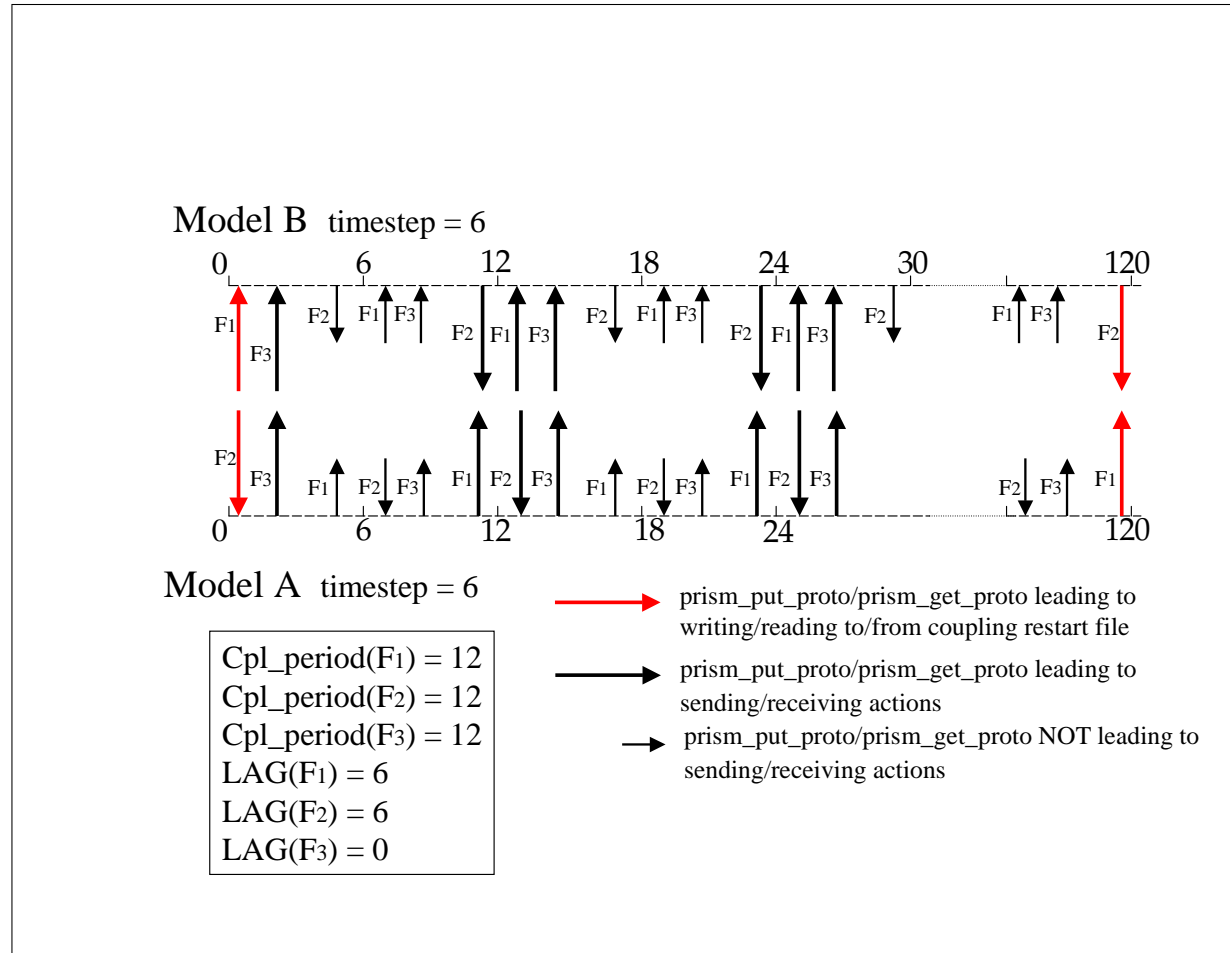


Figure 2.5: LAG concept second example

restart file below the last F_1 oasis_put/prism_put_proto as the date+lag equals the total run time. The analogue is true for F_2 . These values will automatically be read in at the beginning of the next run below the respective oasis_get/prism_get_proto.

2. LAG concept second example

A second coupling algorithm exploiting the LAG concept is illustrated on figure 2.5. During its timestep, model A receives F_2 , sends F_3 and then F_1 ; its timestep length is 6. During its timestep, model B receives F_1 , receives F_3 and then sends F_2 ; its timestep length is also 6. F_1 , F_2 and F_3 coupling periods are both supposed to be equal to 12.

For F_1 and F_2 the situation is similar to the first example. If F_1/F_2 “put” action by model A/B was used at a coupling timestep to match the model B/A “get” action, a deadlock would occur as both models would be waiting on a “get” action. To prevent this, F_1 and F_2 produced at the timestep before have to be used to match the model A and model B “get” actions, which means that a lag of 6 must be defined for both F_1 and F_2 . For both coupling fields, the oasis_put/prism_put_proto performed at times 6 and 18 by the source model then lead to “put” actions (as $6 + 6 = 12$ and $18 + 6 = 24$ which are coupling periods) that match the “get” action performed at time 12 and 24 below the oasis_get/prism_get_proto called by the target model.

For F_3 , sent by model A and received by model B, no lag needs to be defined: the coupling field produced by model A at the coupling timestep can be “consumed” by model B without causing a deadlock situation.

As in the first example, the oasis_get/prism_get_proto performed at the beginning of the run for F_1 and F_2 , will automatically receive data read from their coupling restart files, and the last oasis_put/prism_put_proto performed at the end of the run automatically write them

to their coupling restart file. For F_3 , no coupling restart file is needed nor used as at each coupling period the coupling field produced by model A can be directly “consumed” by model B.

We see here how the introduction of appropriate LAG indices results in “get” in the target model, coupling fields produced by the source model the timestep before; this is, in some coupling configurations, essential to avoid deadlock situations.

2.10.2 The sequence concept

The order of coupling operations in the system is determined solely by the order of calls to send (put) and receive (get) data in the models in conjunction with the setting of the lag in the *namcouple*. Data that is received (get) is always blocking while data that is sent (put) is non-blocking with respect to the model making that call. It is possible to deadlock the system if the relative orders of puts and gets in different models are not compatible.

With OASIS3-MCT, the sequence (SEQ) index in the *namcouple* file now provides the coupling layer with an ability to detect a deadlock before it happens and exit. It does this by tracking the order of get and put calls in models compared to the SEQ specified in the *namcouple*. If there are any inconsistencies, the model will abort gracefully with a useable error message before the system deadlocks. If there are any coupling dependencies in the system, use of the SEQ index is recommended for diagnosis but has no impact on the ultimate solution and is NOT required.

Take the following two examples. In both examples, there are two models, each “put” a field to the other at every coupling period without any lags. In the first case, there is no dependency and each model sends (puts) the data first and then receives the data second at each timestep.

model1	model2
-----	-----
put (fld1)	put (fld2)
get (fld2)	get (fld1)

The put from model1 for fld1 is received by the get in model2 and the put from model2 for fld2 is received by the get in model1. In this case, there is no sequencing dependency and the value of SEQ must be identical (or unset) in the *namcouple* description of the fld1 and fld2 coupling. If SEQ is set to 1 for fld1 and 2 for fld2 in this case, then the model will abort because at runtime, the coupling will detect, in model 2, that fld2 was sent before fld1 was received which is out of sequence as defined by the SEQ settings.

In the next example, there is a dependency in the sequencing.

model1	model2
-----	-----
put (fld1)	get (fld1)
	fld2=g (fld1)
get (fld2)	put (fld2)

In model2, fld2 depends on fld1. If this dependency is known, then there is a benefit in using SEQ=1 for fld1 and SEQ=2 for fld2. At runtime, if the sequencing of both model1 and model2 do not match the above diagram, the model will abort gracefully. For instance, if model2 has the dependency shown above but model1 does not have consistent ordering of the put and get as required by model2, then if SEQ is unused, the model will deadlock and hang. If SEQ is set properly, the coupling layer will detect that the required sequence has not been followed and will exit gracefully with an error message.

Again, the SEQ namecouple setting is only diagnostic and is not required.

Chapter 3

The configuration file *namcouple*

The OASIS3-MCT configuration file *namcouple* contains, below pre-defined keywords, all user's defined information necessary to configure a particular coupled run.

The *namcouple* is a text file with the following characteristics:

- the keywords used to separate the information can appear in any order;
- the number of blanks between two character strings is non-significant;
- all lines beginning with # are ignored and considered as comments.
- **blank lines are not allowed.**

The first part of *namcouple* is devoted to configuration of general parameters such as the number of models involved in the simulation or the number of fields. The second part gathers specific information on each coupling (or I/O) field, e.g. their coupling period, the list of transformations or interpolations to be performed by OASIS3-MCT and associated configuring lines (described in more details in chapter 4), etc.

In OASIS3-MCT, several *namcouple* inputs have been deprecated but, for backwards compatibility, they are still allowed. These inputs will be noted in the following text using the notation “UNUSED” and not fully described. Information below these keywords is obsolete in OASIS3-MCT; it will not be read and will not be used.

In the next sections, a simple *namcouple* example is given and all configuring parameters are described. The additional lines containing the different parameters required for each transformation are described in section 4. A realistic *namcouple* can be found in `oasis3-mct/examples/tutorial/data_oasis3/` directory.

3.1 An example of a simple *namcouple*

The following simple *namcouple* configures a run into which an ocean, an atmosphere and an atmospheric chemistry models are coupled. The ocean provides only the SOSSTSST field to the atmosphere, which in return provides the field CONSFTOT to the ocean. One field (COSENHFL) is exchanged from the atmosphere to the atmospheric chemistry, and one field (SOALBEDO) is read from a file by the ocean.

```
#####  
# First section  
#  
# $SEQMODE  
#  
# $CHANNEL
```

```

#
$NFIELDS
    4
#
$JOBNAME
#
$NBMODEL
    3  ocemod  atmmod  chemod  55  70  99
#
$RUNTIME
    432000
#
$INIDATE
#
$MODINFO
#
$NLOGPRT
    2
#
$CALTYPE
#
#####
# Second section
#
$STRINGS
#
# Field 1
SOSSTSST SISUTESU 1 86400 5 sstoc.nc EXPORTED
182 149 128 64 toce atmo LAG=+14400 SEQ=+1
P 2 P 0
LOCTRANS CHECKIN MAPPING BLASNEW CHECKOUT
#
AVERAGE
INT=1
map_toce_atmo_120315.nc src opt
CONSTANT 273.15
INT=1
#
# Field 2
CONSFTOT SOHEFLDO 6 86400 4 flxat.nc EXPORTED
atmo toce LAG=+14400 SEQ=+2
P 0 P 2
LOCTRANS CHECKIN SCRIPR CHECKOUT
#
ACCUMUL
INT=1
BILINEAR LR SCALAR LATLON 1
INT=1
#
# Field 3
COSENHFL SOSENHFL 37 86400 1 flda3.nc IGNOUT

```

```

atmo    atmo  LAG=+7200
LOTRANS
AVERAGE
#
# Field 4
SOALBEDO SOALBEDO  17  86400  0  SOALBEDO.nc  INPUT
#
#####

```

3.2 First section of *namcouple* file

The first section of *namcouple* uses some predefined keywords prefixed by the \$ sign to locate the related information. The \$ sign must be in the second column. The first ten keywords are described hereafter:

- \$SEQMODE: UNUSED
- \$CHANNEL: UNUSED
- \$NFIELDS: On the line below this keyword is the total number of fields exchanged and described in the second part of the *namcouple*.
- \$JOBNAME: UNUSED
- \$NBMODEL: On the line below this keyword is the number of models running in the given experiment followed by CHARACTER*6 variables giving their names, which must correspond to the name announced by each model when calling `oasis_init_comp` or `prism_init_comp_proto` (second argument, see section 2.2).

Then the user may indicate on the same line the maximum Fortran unit number used by the models. In the example, Fortran units above 55, 70, and 99 are free for respectively the ocean, atmosphere, and atmospheric chemistry models. **In all cases, OASIS3-MCT library assumes, during the initialization phase, that units 1025 and 1026 are free and temporarily uses these units to read the *namcouple* and to write corresponding log messages to file `nout.000000`.** After the initialization phase, OASIS3-MCT will still suppose that units above 1024 are free, unless maximum unit numbers are indicated here in the *namcouple*.

- \$RUNTIME: On the line below this keyword is the total simulated time of the run, expressed in seconds.
- \$INIDATE: UNUSED
- \$MODINFO: UNUSED
- \$NLOGPRT: The line below this keyword refers to the amount of information that will be written to the OASIS3-MCT debug files for each model and processor. The value of \$NLOGPRT may be:
 - 0 : one file `debug.root.xx` is open by the master proces of each model and one file `debug_notroot.xx` is open for all the other processes of each model to write only error information.
 - 1 : one file `debug.root.xx` is open by the master proces of each model to write information equivalent to level 10 (see below); one file `debug_notroot.xx` is open for all the other processes of each model to write error information.
 - 2 : one file `debug.xx.xxxxxx` is open by each process of each model to write normal production diagnostics
 - 5 : as for 2 with in addition some initial debug info
 - 10: as for 5 with in addition the routine calling tree
 - 12: as for 10 with in addition some routine calling notes
 - 15: as for 12 with even more debug diagnostics

- 20: as for 15 with in addition some extra runtime analysis
- 30: full debug information

This value can be changed at runtime with the `oasis_set_debug` or `prism_set_debug` routine (see section 2.9).

- `$CALTYPE`: UNUSED

3.3 Second section of *namcouple* file

The second part of the *namcouple*, starting after the keyword `$STRINGS`, contains coupling information for each coupling (or I/O) field. Its format depends on the field status given by the last entry on the field first line (`EXPORTED`, `IGNOUT` or `INPUT` in the example above). The field may be :

- `AUXILARY`: UNUSED
- `EXPORTED`: exchanged between component models and transformed by OASIS3-MCT
- `EXPOUT`: exchanged, transformed and also written to two debug NetCDF files, one before the sending action in the source model below the `oasis_put` or `prism_put_proto` call (after local transformations `LOCTRANS` and `BLASOLD` if present), and one after the receiving action in the target model below the `prism_get_proto` call (after all transformations). This status should be used when debugging the coupled model only. The name of the debug NetCDF file (one per field) is automatically defined based on the field and component model names.
- `IGNORED`: with OASIS3-MCT, this setting is equivalent to and converted to `EXPORTED`
- `IGNOUT`: with OASIS3-MCT, this setting is equivalent to and converted to `EXPOUT`
- `INPUT`: read in from the input file by the target model below the `oasis_get` or `prism_get_proto` call at appropriate times corresponding to the input period indicated by the user in the *namcouple*. See section 5.5 for the format of the input file.
- `OUTPUT`: written out to an output debug NetCDF file by the source model below the `oasis_put` or `prism_put_proto` call, after local transformations `LOCTRANS` and `BLASOLD`, at appropriate times corresponding to the output period indicated by the user in the *namcouple*.

3.3.1 Second section of *namcouple* for `EXPORTED` and `EXPOUT` fields

The first 3 lines for fields with status `EXPORTED` and `EXPOUT` are as follows:

```
SOSSTSST SISUTESU 1 86400 5 sstoc.nc sstat.nc EXPORTED
182 149 128 64 toce atmo LAG=+14400 SEQ=+1
P 2 P 0
```

where the different entries are:

- Field first line:
 - `SOSSTSST` : symbolic name for the field in the source model (CHARACTER*8). It has to match the argument name of the corresponding field declaration in the source model; see `oasis_def_var` or `prism_def_var_proto` in section 2.5.
 - `SISUTESU` : symbolic name for the field in the target model (CHARACTER*8). It has to match the argument name of the corresponding field declaration in the target model; see `oasis_def_var` or `prism_def_var_proto` in section 2.5.
 - `1` : UNUSED but still required for parsing
 - `86400` : coupling and/or I/O period for the field, in seconds.
 - `5` : number of transformations to be performed by OASIS3 on this field.

- sstoc.nc : name of the coupling restart file for the field (CHARACTER*8); mandatory even if no coupling restart file is effectively used. (for more detail, see section 5.3);
- sstat.nc : UNUSED but still required for parsing
- EXPORTED : field status.
- Field second line:
 - 182 : number of points for the source grid first dimension (optional)¹.
 - 149 : number of points for the source grid second dimension (optional)¹.
 - 128 : number of points for the target grid first dimension (optional)¹.
 - 64 : number of points for the target grid second dimension (optional)¹.
 - toce : prefix of the source grid name in grid data files (see section 5.2) (CHARACTER*4)
 - atmo : prefix of the target grid name in grid data files (CHARACTER*4)
 - LAG=+14400: optional lag index for the field expressed in seconds
 - SEQ=+1: optional sequence index for the field (see section 2.10)
- Field third line
 - P : source grid first dimension characteristic ('P': periodical; 'R': regional).
 - 2 : source grid first dimension number of overlapping grid points.
 - P : target grid first dimension characteristic ('P': periodical; 'R': regional).
 - 0 : target grid first dimension number of overlapping grid points.

The fourth line gives the list of transformations to be performed for this field. In addition, there is one or more configuring lines describing some parameters for each transformation. These additional lines are described in more details in the chapter 4.

3.3.2 Second section of *namcouple* for OUTPUT fields

The first 2 lines for fields with status OUTPUT are as follows:

```
COSHFTOT COSHFTOT 7 86400 0 fldhftot.nc OUTPUT
atmo atmo
```

where the different entries are as for EXPOUT fields, except that the source symbolic name must be repeated twice on the field first line, the restart file name is needed only if LOCTRANS transformations are present, there is no output file name on the first line and no LAG or SEQ index at the end of the second line. The name of the output file is automatically defined based on the field and component model names. The third line is LOCTRANS if this transformation is chosen for the field. Note that LOCTRANS is the only transformation supported for OUTPUT fields.

3.3.3 Second section of *namcouple* for INPUT fields

The first and only line for fields with status INPUT is:

```
SOALBEDO SOALBEDO 17 86400 0 SOALBEDO.nc INPUT
```

- SOALBEDO: symbolic name for the field in the target model (CHARACTER*8 repeated twice)
- 17: index in auxiliary file cf_name_table.txt (see above for EXPORTED fields)
- 86400: input period in seconds
- 0: number of transformations (always 0 for INPUT fields)

¹The 2D dimensions of the grids must be provided in the *namcouple* so to have 2D fields in the debug files; otherwise, the fields in the debug files will be 1D.

-
- `SOALBEDO.nc`: CHARACTER*32 giving the input file name (for more detail on its format, see section 5.5)
 - `INPUT`: field status.

Chapter 4

Transformations and interpolations

Different transformations and 2D interpolations are available in OASIS3-MCT to adapt the coupling fields from a source model grid to a target model grid. In the following paragraphs, a description of each transformation with its corresponding configuring lines is given. Features that are now deprecated (non functional) compared to prior versions will be noted with the string `UNUSED` but not described.

4.1 Time transformations

- **LOCTRANS:**

LOCTRANS requires one configuring line on which a time transformation, automatically performed below the call to `oasis_put` or `prism_put_proto`, should be indicated:

```
# LOCTRANS operation
$TRANSFORM
```

where `$TRANSFORM` can be

- **INSTANT:** no time transformation, the instantaneous field is transferred;
- **ACCUMUL:** the field accumulated over the previous coupling period is exchanged (the accumulation is simply done over the arrays `field_array` provided as third argument to the `oasis_put` or `prism_put_proto` calls, not weighted by the time interval between these calls);
- **AVERAGE:** the field averaged over the previous coupling period is transferred (the average is simply done over the arrays `field_array` provided as third argument to the `oasis_put` or `prism_put_proto` calls, not weighted by the time interval between these calls);
- **T_MIN:** the minimum value of the field for each source grid point over the previous coupling period is transferred;
- **T_MAX:** the maximum value of the field for each source grid point over the previous coupling period is transferred;
- **ONCE:** only one `oasis_put / prism_put_proto` or `oasis_get / prism_get_proto` will be performed; this is equivalent to giving the length of the run as coupling or I/O period.

With OASIS3-MCT, time transformations are supported more generally with use of the coupling restart file. The coupling restart file allows the partial time transformation to be saved at the end of a run for exact restart at the start of the next run. For that reason, coupling restart filenames are now required for all *namcouple* transformations that use LOCTRANS (with non INSTANT values). In this mode, OASIS3-MCT will exit gracefully with an error message if a restart filename is not

provided. This is the reason an optional restart file is now provided on the OUTPUT *namcouple* input line.

4.2 The pre-processing transformations

- **REDGLO** UNUSED
- **INVERT**: UNUSED
- **MASK**: UNUSED
- **EXTRAP**: UNUSED
- **CHECKIN**:

CHECKIN calculates the global minimum, the maximum and the sum of the the source field values and prints them to the debug file (for the master process of the source component model only). This operation does not transform the field.

The generic input line is as follows, even if \$NINT has no impact in OASIS3-MCT:

```
# CHECKIN operation
$INT = $NINT
```

- **CORRECT**: UNUSED

4.3 The interpolation

- **BLASOLD**:

BLASOLD allows the source field to be scaled and allows a scalar to be added to the field. The prior ability to perform a linear combination of the current coupling field with other coupling fields has been deprecated in OASIS3-MCT. This transformation occurs before the interpolation *per se*.

This transformation requires at least one configuring line with two parameters:

```
# BLASOLD operation
$XMULT $NBFIELDS
```

where \$XMULT is the multiplicative coefficient of the source field. \$NBFIELDS must be 0 if no scalar needs to be added or 1 if a scalar needs to be added. In this last case, an additional input line is required where \$AVALUE is the scalar to be added to the field :

```
CONSTANT $AVALUE
```

- **MAPPING**:

The **MAPPING** keyword is used to specify an input file to be read and used for mapping (ie. regridding or interpolation); the **MAPPING** file must follow the **SCRIPR** format. This is an alternative method to **SCRIPR** for setting the mapping file.

In the current implementation, each pair of source and target points in the **MAPPING** file can be linked by only one weight, i.e. remappings such as **SCRIPR/BICUBIC** involving at each source grid point the value of the field, of the gradients and the cross-gradient, or second-order conservative remapping are not supported.

This transformation requires at least one configuring line with one filename and two optional string values:

```
$MAPNAME $MAPLOC $MAPSTRATEGY
```

- \$MAPNAME is the name of the mapping file to read. This is a netcdf file consistent with the **SCRIPR** map file format (see section 5.5).

- `$MAPLOC` is optional and can be either `src` or `dst`. With `src`, the mapping will be done in parallel on the source processors before communication to the destination model and processors; this is the default. With `dst`, the mapping is done on the destination processors after the data is sent from the source model on the source grid.
- `$MAPSTRATEGY` is optional and can be either `bfb`, `sum`, or `opt`. In `bfb` mode, the mapping is done using a strategy that produces bit-for-bit identical results regardless of the grid decompositions without leveraging a partial sum computation. With `sum`, the transform is done using the partial sum approach which generally introduces roundoff level changes in the results on different processor counts. Option `opt` allows the coupling layer to choose either approach based on an analysis of which strategy is likely to run faster. Usually, partial sums will be used if the source grid has a higher resolution than the target grid as this should reduce the overall communication.

Note that if `SCRIPR` (see below) is used to calculate the remapping file, `MAPPING` can still be listed in the `namcouple` to specify a name for the remapping file generated by `SCRIPR` different from the default and/or to specify a `$MAPLOC` or `$MAPSTRATEGY` option.

• **SCRIPR:**

`SCRIPR` gathers the interpolation techniques offered by Los Alamos National Laboratory `SCRIP` 1.4 library (Jones 1999)¹. `SCRIPR` routines are in `oasis3-mct/lib/scrip`. See the `SCRIP` 1.4 documentation in `oasis3/doc/SCRIPusers.pdf` for more details on the interpolation algorithms. In the current implementation, each pair of source and target points in the `MAPPING` file can be linked by only one weight, i.e. remappings such as `SCRIPR/BICUBIC` involving at each source grid point the value of the field, of the gradients and the cross-gradient, or second-order conservative `SCRIPR/CONSERV/SECOND` remapping are not supported.

The following types of interpolations are available: `ICICICI`

- `DISTWGT` performs a distance weighted nearest-neighbour interpolation (`N` neighbours). All types of grids are supported.
 - * Masked target grid points: the zero value is associated to masked target grid points.
 - * Non-masked target grid points having some of the `N` source nearest neighbours masked: a nearest neighbour algorithm using the remaining non masked source nearest neighbours is applied.
 - * Non-masked target grid points having all of the `N` source nearest neighbours masked: by default, the nearest non-masked source neighbour is used.

The configuring line is:

```
# SCRIPR (for DISWGT)
    $CMETH $CGRS $CFTYP $REST $NBIN $NV $ASSCMP $PROJCART

* $CMETH = DISTWGT.
* $CGRS is the source grid type (LR, D or U)- see annexe A.
* $CFTYP is the field type: SCALAR if the field is a scalar one, or VECTOR_I or VECTOR_J
  whether the field represents respectively the first or the second component of a vector field
  (see paragraph Support of vector fields below). The option VECTOR, which in fact leads
  to a scalar treatment of the field (as in the previous versions), is still accepted.
* $REST is the search restriction type: LATLON or LATITUDE (see SCRIP 1.4 documenta-
  tion SCRIPusers.pdf). Note that for D or U grid, the restriction may influence slightly
  the result near the borders of the restriction bins. (XXX to be checked)
* $NBIN the number of restriction bins (see SCRIP 1.4 documentation SCRIPusers.pdf).
```

¹See also <http://climate.lanl.gov/Software/SCRIP/> and the copyright statement in appendix 1.3.3.

- * \$NV is the number of neighbours used.
 - * \$ASSCMP: optional, for VECTOR_I or VECTOR_J vector fields only; the source symbolic name of the associated vector component.
 - * \$PROJCART: optional, for vector fields only; should be PROJ CART if the user wants the vector components to be projected in a Cartesian coordinate system before interpolation (see paragraph **Support of vector fields** below).
- GAUSWGT performs a N nearest-neighbour interpolation weighted by their distance and a gaussian function. All grid types are supported.
- * Masked target grid points: the zero value is associated to masked target grid points.
 - * Non-masked target grid points having some of the N source nearest neighbours masked: a nearest neighbour algorithm using the remaining non masked source nearest neighbours is applied.
 - * Non-masked target grid points having their N nearest neighbours all masked: the zero value will be associated to these target points.

The configuring line is:

```
# SCRIPR (for GAUSWGT)
    $CMETH $CGRS $CFTYP $REST $NBIN $NV $VAR $ASSCMP $PROJCART
```

all entries are as for DISTWGT, except that:

- * \$CMETH = GAUSWGT
 - * \$VAR, which must be given as a REAL value (e.g 2.0 and not 2), defines the weight given to a neighbour source grid point as proportional to $\exp(-1/2 \cdot d^2/\sigma^2)$ where d is the distance between the source and target grid points, and $\sigma^2 = \$VAR \cdot \bar{d}^2$ where \bar{d}^2 is the average distance between two source grid points (calculated automatically by OASIS3).
- BILINEAR performs an interpolation based on a local bilinear approximation (see details in chapter 4 of SCRIP 1.4 documentation SCRIPusers.pdf)

For BILINEAR and BICUBIC, Logically-Rectangular (LR) and Reduced (D) source grid types are supported.

- * Masked target grid points: the zero value is associated to masked target grid points.
- * Non-masked target grid points having some of the source points normally used in the bilinear or bicubic interpolation masked: a N nearest neighbour algorithm using the remaining non masked source points is applied.
- * Non-masked target grid points having their N nearest neighbours all masked: the zero value will be associated to these target points.

The configuring line is:

```
# SCRIPR (for BILINEAR or BICUBIC)
    $CMETH $CGRS $CFTYP $REST $NBIN $ASSCMP $PROJCART
```

- * \$CMETH = BILINEAR or BICUBIC
 - * \$CGRS is the source grid type (LR or D)
 - * \$CFTYP, \$NBIN, \$ASSCMP \$PROJCART are as for DISTWGT.
 - * \$REST is as for DISTWGT, except that only LATITUDE is possible for a Reduced (D) source grid.
- BICUBIC is currently unsupported as a mapping option because it is higher order. It performs an interpolation based on a local bicubic approximation (see details in chapter 5 of SCRIP 1.4 documentation SCRIPusers.pdf) See BILINEAR for configure line

- CONSERV performs 1st or 2nd order conservative remapping, which means that the weight of a source cell is proportional to area intersected by the target cell. Note that 2nd order conservative mapping files can be generated but not used currently.

The configuring line is:

```
# SCRIPR (for CONSERV)
    $CMETH $CGRS $CFTYP $REST $NBIN $NORM $ORDER $ASSCMP $PROJCART

* $CMETH = CONSERV
* $CGRS is the source grid type: LR, D and U are supported for first-order remapping
  if the grid corners are given by the user in the grid data file grids.nc ; only LR is
  supported if the grid corners are not available in grids.nc and therefore have to be
  calculated automatically by OASIS3. For second-order remapping, only LR is supported
  because the gradient of the coupling field used in the transformation has to be calculated
  automatically by OASIS3.
* $CFTYP, $REST, $NBIN, $ASSCMP, and $PROJCART are as for DISTWGT.
* $NORM is the NORMalization option:
  · FRACAREA: The sum of the non-masked source cell intersected areas is used to
    NORMalise each target cell field value: the flux is not locally conserved, but the
    flux value itself is reasonable.
  · DESTAREA: The total target cell area is used to NORMalise each target cell field
    value even if it only partly intersects non-masked source grid cells: local flux conser-
    vation is ensured, but unreasonable flux values may result.
  · FRACNNEI: as FRACAREA, except that at least the source nearest unmasked neigh-
    bour is used for unmasked target cells that intersect only masked source cells. Note
    that a zero value will be assigned to a target cell that does not intersect any source
    cells (masked or unmasked), even with FRACNNEI option.
* $ORDER: FIRST or SECOND for first or second order remapping respectively (see SCRIP
  1.4 documentation). Note that CONSERV/SECOND is not positive definite and has not
  been fully validated yet.
```

Precautions related to the use of the SCRIPR/CONSERV remapping in particular

- For the 1st order conservative remapping: the weight of a source cell is proportional to area of the source cell intersected by target cell. Using the divergence theorem, the SCRIP library evaluates this area with the line integral along the cell borders enclosing the area. As the real shape of the borders is not known (only the location of the 4 corners of each cell is known), the library assumes that the borders are linear in latitude and longitude between two corners. In general, this assumption becomes less valid closer to the pole and for latitudes above the `north_thresh` or below the `south_thresh` values specified in `oasis3/lib/scrip/remap-conserv.F`, the library evaluates the intersection between two border segments using a Lambert equivalent azimuthal projection. Problems have been observed in some cases for the grid cell located around this `north_thresh` or `south_thresh` latitude.
- Another limitation of the SCRIP 1st order conservative remapping algorithm is that it also supposes, for line integral calculation, that $\sin(\text{latitude})$ is linear in longitude on the cell borders which again is in general not valid close to the pole.
- For a proper conservative remapping, the corners of a cell have to coincide with the corners of its neighbour cell.
- If two cells of a grid overlay, at least the one with the greater numerical index must be masked (they also can be both masked) for a proper conservative remapping. For example, if the grid

line with $i=1$ overlaps the grid line with $i=\text{imax}$, it is the latter that must be masked. When this is not the case with the mask defined in *masks.nc*, OASIS must be compiled with the CPP key `TREAT_OVERLAY` which will ensure that these rules are respected. This CPP key was introduced in *oasis3_3* version.

- A target grid cell intersecting no source cell (either masked or non masked) at all i.e. falling in a “hole” of the source grid will in all cases get a zero value.
- If a target grid cell intersects only masked source cells, it will still get a zero value unless:
 - the `FRACNNEI` normalisation option is used, in which case it will get the nearest non masked neighbour value, or
 - the routines *oasis3/lib/scrip/src/scriprmp.f* or *vector.F90* - for vector interpolation - are compiled with `ll_weightot=.true.` in which case, the value `1.0E+20` will be assigned to these target grid cell intersecting only masked source cells (for easier identification).

Precautions related to the use of the SCRIPR remappings in general

- For using SCRIPR interpolations, linking with the NetCDF library is mandatory and the grid data files (see section 5.2) must be NetCDF files (binary files are not supported).
- When the SCRIP library performs a remapping, it first checks if the file containing the corresponding remapping weights and addresses exists. If it exists, it reads them from the file; if not, it calculates them and store them in a file. The file is created in the working directory and is called *rmp_src_tg_tgt_INTTYPE_NORMAOPT.nc*, where *srcg* and *tggt* are the acronyms of respectively the source and the target grids, *INTTYPE* is the interpolation type (i.e. `DISTWGT`, `GAUSWGT`, `BILINEA`, `BICUBIC`, or `CONSERV`) and *NORMAOPT* is the normalization option (i.e. `DESTAREA`, `FRACAREA` or `FRACNNEI` for `CONSERV` only). The problem comes from the fact that the weights and addresses will also differ whether or not the `MASK` and `EXTRAP` transformations are first activated during the pre-processing phase (see section 4.2) and this option is not stored in the remapping file name. Therefore, the remapping file used will be the one created for the first field having the same source grid, target grid, and interpolation type (and the same normalization option for `CONSERV`), even if the `MASK` and `EXTRAP` transformations are used or not for that field. (This inconsistency is however usually not a problem as the `MASK` and `EXTRAP` transformations are usually used for all fields having the same source grid, target grid, and interpolation type, or not at all.)

Support of vector fields with the SCRIPR remappings

Vector mapping is NOT supported in this version of OASIS. Vector fields are treated as if each field were an independent scalar field. Vector mapping capabilities will be added in future versions.

- **INTERP:** UNUSED
- **MOZAIC:** UNUSED
- **FILLING:** UNUSED

4.4 The “cooking” stage

The following transformations are available in the “cooking” part of OASIS3.

- **CONSERV:**

`CONSERV` ensures a global modification of the coupling field. This analysis requires the areas of the source and target grid meshes to be transferred to the coupler with the *oasis_write_area/prism_write_area* call (see section 2.3). In the *namcouple*, `CONSERV` requires one input line with one argument and one optional argument:

```
# CONSERV operation
    $CMETH $CONSOPT
```

where:

- \$CMETH is the method desired with the following choices
 - * with \$CMETH = GLOBAL, the field is integrated on both source and target grids, without considering values of masked points, and the residual (target - source) is uniformly distributed on the target grid; this option ensures global conservation of the field
 - * with \$CMETH = GLBPOS, the same operation is performed except that the residual is distributed proportionally to the value of the original field; this option ensures the global conservation of the field and does not change the sign of the field
 - * with \$CMETH = BASBAL, the operation is analogous to GLOBAL except that the non masked surface of the source and the target grids are taken into account in the calculation of the residual; this option does not ensure global conservation of the field but ensures that the energy received is proportional to the non masked surface of the target grid
 - * with \$CMETH = BASPOS, the non masked surface of the source and the target grids are taken into account and the residual is distributed proportionally to the value of the original field; therefore, this option does not ensure global conservation of the field but ensures that the energy received is proportional to the non masked surface of the target grid and it does not change the sign of the field.
- \$CONSOPT is an optional argument specifying the algorithm. \$CONSOPT can be `bfb` or `opt`. The `bfb` option enforces a bit-for-bit transformation regardless of the grid decomposition or process count. The `opt` option carries out the conservation using an optimal algorithm using less memory and a faster approach. Option `bfb` is the default setting.

Note that for this operation to be correct, overlapping grid cells on the source grid or on the target grid must be masked.

- **SUBGRID:** UNUSED
- **BLASNEW:**

BLASNEW performs a scalar multiply or scalar add to any destination field. This is the equivalent of BLASOLD on the destination side. The prior feature that supported linear combinations of the current coupling field with any other fields after the interpolation has been deprecated.

This analysis requires the same input line as BLASOLD.

- **MASKP:** UNUSED

4.5 The post-processing

The following analyses are available in the post-processing part of OASIS3.

- **REVERSE:** UNUSED
- **CHECKOUT:**

CHECKOUT calculates the global minimum, the maximum and the sum of the the target field values and prints them to the OASIS debug file (for the master process of the target component model only). This operation does not transform the field. The generic input line is as for CHECKIN (see above).

- **GLORED:** UNUSED

Chapter 5

OASIS3 auxiliary data files

OASIS3 needs auxiliary data files describing coupling and I/O field names and units, defining the grids of the models being coupled, containing the field coupling restart values or input data values, as well as a number of other auxiliary data files used in specific transformations.

5.1 Field names and units

The text file `cf_name_table.txt`, has been available to users in the past. Use of this file is currently not supported but will be in future versions.

5.2 Grid data files

The grids of the models being coupled must be given by the user, or directly by the model through PSMILe specific calls that write grid data. Note that if the grid data files exist in the working directory, fields in those files will not be overwritten by the PSMILe specific calls (see section 2.3). These files are netCDF.

The arrays containing the grid information are dimensioned (nx, ny) , where nx and ny are the grid first and second dimension. Unstructured grids are supported by setting the ny dimension to 1 and then nx is the total number of grid points.

1. *grids.nc*: contains the model grid longitudes, latitudes, and local angles (if any) in single or double precision REAL arrays (depending on OASIS3 compilation options). The array names must be composed of a prefix (4 characters), given by the user in the *namcouple* on the second line of each field (see section 3.3), and of a suffix (4 characters); this suffix is “.lon” or “.lat” for respectively the grid point longitudes or latitudes.

If the SCRIPR/CONSERV remapping is used, longitudes and latitudes for the source and target grid **corners** must also be available in the *grids.nc* file as arrays dimensioned $(nx, ny, 4)$ or $(nbr_pts, 1, 4)$ where 4 is the number of corners (in the counterclockwise sense). The names of the arrays must be composed of the grid prefix and the suffix “.clo” or “.cla” for respectively the grid corner longitudes or latitudes. As for the other grid information, the corners can be provided in *grids.nc* before the run by the user or directly by the model through PSMILe specific calls (see section 2.3).

Longitudes must be given in degrees East in the interval -360.0 to 720.0. Latitudes must be given in degrees North in the interval -90.0 to 90.0. Note that if some grid points overlap, it is recommended to define those points with the same number (e.g. 360.0 for both, not 450.0 for one and 90.0 for the other) to ensure automatic detection of overlap by OASIS.

The corners of a cell cannot be defined modulo 360 degrees. For example, a cell located over Greenwich will have to be defined with corners at -1.0 deg and 1.0 deg but not with corners at 359.0

deg and 1.0 deg.

Cells larger than 180.0 degrees in longitude are not supported.

If vector fields are defined on a grid which has a local coordinate system not oriented in the usual zonal and meridional directions, the local angle of the grid coordinate system must be given in *grids.nc* file in an array which name must be composed of the grid prefix and the suffix “.ang”. The angle is defined as the angle between the first component and the zonal direction (which is also the angle between the second component and the meridional direction). If one of the `SCRIPR` interpolations is requested for a vector field, OASIS3 automatically performs the rotation from the local coordinate system to the geographic spherical coordinate system for a source grid, or vice-versa for a target grid.

2. *masks.nc*: contains the masks for all component model grids in `INTEGER` arrays (0 -not masked i.e. active- or 1 -masked i.e. not active- for each grid point). The array names must be composed of the grid prefix and the suffix “.msk”. This file, *masks* or *masks.nc*, is mandatory.
3. *areas.nc*: this file contains mesh surfaces for the component model grids in single or double precision `REAL` arrays (depending on OASIS3 compilation options). The array names must be composed of the grid prefix and the suffix “.srf”. The surfaces may be given in any units but they must be all the same (in `INTERP/GAUSSIAN`, it is assumed that the units are m^2 but they are used for statistics calculations only.) This file *areas* or *areas.nc* is mandatory for `CHECKIN`, `CHECKOUT` or `CONSERV`, and used for statistic calculations in `INTERP/GAUSSIAN`; it is not required otherwise.

5.3 Coupling restart files

At the beginning of a coupled run, some coupling fields may have to be initially read from their coupling restart file on their source grid (see section 2.10). When needed, these files are also automatically updated by the last `prismput_proto` call of the run (see section 2.7.1) . **Warning:** the date is not written or read to/from the restart file; therefore, the user has to make sure that the appropriate restart file is present in the working directory.

Note that all restart files have to be present in the working directory at the beginning of the run even if one model is delayed with respect to the others.

The name of the coupling restart file is given by the 6th character string on the first configuring line for each field in the *namcouple* (see section 3.3). Coupling fields coming from different models cannot be in the same coupling restart files, but for each model, there can be an arbitrary number of fields written in one coupling restart file.

In the coupling restart files, the fields must be provided on the source grid in single or double precision `REAL` arrays and, as the grid data files, must be dimensioned (nx, ny) , where nx and ny are the grid first and second dimension, except for fields given on unstructured for which the arrays are dimensioned $(nt, 1)$, where nt is the total number of grid points. The shape and orientation of each restart field (and of the corresponding coupling fields exchanged during the simulation) must be coherent with the shape of its grid data arrays.

NetCDF formats are supported; for NetCDF file the suffix *.nc* is not mandatory. In the NetCDF restart files, the field arrays must have the source symbolic name indicated in the *namcouple*. (see section 3.3).

5.4 Mapping files

The mapping files to be read using the `MAPPING` option are consistent with the files generated in OASIS3. The files are *netcdf* and the key fields are *num_links* (the number of weights for each grid pair), *src_grid_size* (the 2d size of the source grid), *dst_grid_size* (the 2d size of the destination grid), *num_wgts*

(the total number of weights), `remap_matrix` (the weights), `dst_address` (the global destination index for the weight), `src_address` (the global source index for the weight).

5.5 Input data files

Fields with status `INPUT` in the *namcouple* will, at runtime, simply be read in from a NetCDF input file by the target model PSMILe below the `prism_get_proto` call, at appropriate times corresponding to the input period indicated by the user in the *namcouple*.

The name of the file must be the one given on the field first configuring line in the *namcouple* (see section 3.3.3). There must be one input file per `INPUT` field, containing a time sequence of the field in a single or double precision REAL array named with the field symbolic name in the *namcouple* and dimensioned `(nx, ny, time)` or `(nbr_pts, 1, time)`. The time variable has to be an array `time(time)` expressed in “seconds since beginning of run”. The “time” dimension has to be the unlimited dimension.

5.6 Transformation auxiliary data files

Some transformations need auxiliary data files. In particular, interpolation requires that mapping files be generated either offline and set through the `MAPPING namcouple` keyword or be generated via scrip in OASIS using the `SCRIPR namcouple` keyword.

5.6.1 Auxiliary data files for SCRIPR

The NetCDF files containing the weights and addresses for the `SCRIPR` remappings (see section 4.3) are automatically generated at runtime by OASIS3. Their structure is described in detail in section 2.2.3 of the `SCRIP` documentation available in `oasis3/doc/SCRIPusers.pdf`. In particular, they are netCDF files with the following one fields

- `src_grid_size` is a scalar integer indicating the total number of global grid cells for the source grid. This field in a netCDF dimension.
- `dst_grid_size` is a scalar integer indicating the total number of global grid cells for the destination (or target) grid. This field in a netCDF dimension.
- `num_links` is a scalar integer indicating the total number of associated grid pairs in the file. This is typically a large number. This field is a netCDF dimension.
- `num_wgts` is a scalar integer indicating the number of weights per associated grid pair. For first order mapping, this is 1. This field in a netCDF dimension.
- `src_address` is a one dimensional array of size `num_links`. It contains the integer source address associated with each weight. This field is a netCDF variable.
- `dst_address` is a one dimensional array of size `num_links`. It contains the integer destination address associated with each weight. This field is a netCDF variable.
- `remap_matrix` is a two dimensional array of size `num_links, num_wgts`. It contains the real weight value(s) associated with the source and destination address. This field is a netCDF variable.

The mapping implementation does the following using a parallel approach

```
! this is pseudocode for first order mapping

! src_arr is an input, dst_array is an output
! src_address, dst_address, remap_matrix are read from mapping file
! while remap_matrix supports num_wgts > 1, for first order
!   num_wgts = 1 and the sparse matrix multiply below is hardcoded
```

```
!   for first order mapping only.

real :: src_arr(:), dst_arr(:), remap_matrix(:, :)
integer :: src_address(:), dst_address(:)

allocate(src_arr(src_grid_size))
allocate(dst_arr(dst_grid_size))

allocate(src_address(num_links))
allocate(dst_address(num_links))
allocate(remap_matrix(num_wgts, num_links))

read(src_address)
read(dst_address)
read(remap_matrix)

src_arr(:) = from_input(:)
dst_arr(:) = 0.0
do n = 1, num_links
    sindex = src_address(n)
    dindex = dst_address(n)
    dst_arr(dindex) = dst_arr(dindex) + src_array(sindex)*remap_matrix(1,n)
enddo
```

Chapter 6

Compiling and running OASIS3 and TOYOASIS3

6.1 Compiling OASIS3 and debugging

6.1.1 Compilation with TopMakefileOasis3

Compiling OASIS3 can be done in directory `oasis3/util/make_dir` with Makefile `TopMakefileOasis3` which must be completed with a header file `make.your_platform` specific to the compiling platform used and specified in `oasis3/util/make_dir/make.inc`. One of the header files distributed with the release can be used as a template. The root of the OASIS3 tree can be anywhere and must be set in the variable `COUPLE` in the `make.your_platform` file.

The following commands are available:

- `make -f TopMakefileOasis3`
compiles OASIS3 libraries *mct* and *scrip*
- `make realclean -f TopMakefileOasis3:`
removes OASIS3 and PSMILe library compiled sources and librairies.

Log and error messages from compilation are saved in the files `COMP.log` and `COMP.err` in `make_dir`.

During compilation, a new compiling directory, defined by variable `ARCHDIR` is created. After successful compilation, resulting executables are found in the compiling directory in `/bin`, libraries in `/lib` and object and module files in `/build`.

The different pre-compiling flags used for OASIS3 and its associated PSMILe library are described in section 6.1.2.

6.1.2 CPP keys

The following CPP keys are coded in OASIS3 and associated PSMILe library and can be specified in `CPPDEF` in `make.your_platform` file.

- To ensure, in `SCRIPR/CONSERV` remapping (see section 4.3), that if two cells of the source grid overlay, at least the one with the greater numerical index is masked (they also can be both masked); this is mandatory for this remapping. For example, if the grid line with `i=1` overlaps the grid line with `i=imax`, it is the latter that must be masked; when this is not the case with the mask defined in *masks.nc*, this CPP key forces these rules are to be respected.

– `TREAT_OVERLAY`

- To reproduce default behaviour of SCRIPR/DISTWGT before version oasis3_3, i.e. the zero value is associated to the target points having all of the N source nearest neighbours masked (see section 4.3 for details).
 - NOT_NNEIGHBOUR

6.1.3 Debugging

If you experience problems while running your coupled model with OASIS3, you can obtain more information on what is happening with:

- Increasing \$NLOGPRT in your *namcouple* (see section 3.2)

6.2 Running OASIS3 in coupled mode with TOYOASIS3

In order to test the OASIS3 coupler in a light coupled configuration, CERFACS has written 3 “toy” component models, mimicking an atmosphere model (atmoa3), an ocean model (oceoa3), and a chemistry model (cheoa3), which sources can be found in `oasis3/examples/toyoasis3/src`. These “toy” component models are ‘empty’ in the sense that they do not model any real physics or dynamics. The coupled combination of these 3 “toy” component models through OASIS3 coupling software is referred to as the TOYOASIS3 coupled model; the TOYOASIS3 coupling is realistic as the coupling algorithm linking the toy component models, the size and the grid of the 2D coupling fields, and the operations performed by OASIS3 on the coupling fields are realistic.

The current version of OASIS3 and its TOYOASIS3 example coupled model was successfully compiled and run on:

- Linux neolith1 2.6.18-194.8.1.el5, with ifort/icc and openmpi - 1.4 and scalimpi 3.13.8, thanks to C. Basu from NSC (Sweden)
- cluster BULLX based on Intel/westmere, with Intel 11.1.073 compiler and “bullxmpi 1.0.2” MPI library (compilation only)
- CRAY XT platforms, with Cray Fortran Compiler crayftn and MPI library xt-mpt, thanks to C. Henriot from Cray
- (XXX to be completed)

Previous versions were compiled and run on many other platforms.

Compiling OASIS3 was described in section 6.1. In the following section, the TOYOASIS3 example coupled model is first described in more detail (see section 6.2.1), then instructions on how to compile and run TOYOASIS3 are given in section 6.2.2.

6.2.1 TOYOASIS3 description

The oceoa3 model

The oceoa3 model has a 2D logically-rectangular, stretched and rotated grid of 182x149 points, which corresponds to a real ocean model grid (with two poles of convergence over America and Asia). Oceoa3 timestep is 14400 seconds; it performs therefore 36 timesteps per 6-day run.

OASIS3 PSMILe routines are detailed in section 2. At the beginning of a run, oceoa3 performs appropriate PSMILe calls to initialize the coupling, define its grids, and declare its I/O or coupling fields. As oceoa3 is not parallel, it calls the PSMILe `prism_def_partition` routine to define only one Serial partition containing the 182X149 grid points.

Then, oceoa3 starts its timestep loop. At the beginning of its timestep, oceoa3 calls the PSMILe `prism_get` routine 6 times to request the fields named Field3, Field4, Field6 to Field9 on table 6.1. At the end of its

timestep, oceoa3 calls PSMILe prism_put routine to send fields named Field1 and Field2 on table 6.1. The fields will be effectively received or sent only at the coupling frequency defined by the user (see section 3.3).

Finally, at the end of the run, oceoa3 performs the PSMILe finalization call.

The atmoa3 model

The atmoa3 model has a realistic atmospheric Gaussian reduced grid with 6232 points. Its timestep is 3600 seconds; it therefore performs 144 timesteps per 6-day run.

As oceoa3, atmoa3 performs, at the beginning of a run, appropriate PSMILe calls to initialize the coupling, define its grids, and declare its I/O or coupling variables. Then atmoa3 retrieves a local communicator for its internal parallelization with a call to PSMILe prism_get_localcomm routine, useful if the MPI1 communication technique is chosen by the user (see section 2.2), and defines its local partition calling the PSMILe prism_def_partition routine.

Then, atmoa3 starts its timestep loop. At the beginning of its timestep, atmoa3 calls the PSMILe prism_get routine 3 times to request the fields named Field1, Field2 and Field11 on table 6.1. At the end of its timestep, atmoa3 calls PSMILe prism_put routine to send fields named Field4 to Field10 on table 6.1. The fields will be effectively received or sent only at the coupling frequency defined in the *namcouple* (see section 3.3) of the coupled model that one can find in *oasis3/examples/toyoasis3/input*.

Finally, at the end of the run, atmoa3 performs the PSMILe finalization call.

The cheoa3 model

Cheoa3 is integrated on the same atmospheric model grid than atmoa3. Its timestep is 7200 seconds; it therefore performs 72 timesteps per 6-day run.

As the other toymodels, cheoa3 performs, at the beginning of a run, appropriate PSMILe calls to initialize the coupling, define its grids, and declare its I/O or coupling variables; it also retrieves a local communicator if needed. As cheoa3 has the same grid than atmoa3, a direct exchange of coupling fields can occur between those two models, without going through OASIS3 interpolation process. To insure this, the coupling field must have a field status 'IGNORED' or 'IGNOUT' in the OASIS3 configuration file *namcouple* (see section 3.3) and the two models must have also the same parallel decomposition. Cheoa3 decomposition is hardcoded the same way than atmoa3, and if the user modifies the atmoa3 decomposition, he has to modify the cheoa3 decomposition the same way by changing cheoa3 values for *il_nbcplproc* and *cdec* (see below).

At the beginning of its timestep, cheoa3 calls the PSMILe prism_get routine to request Field10 (see table 6.1). At the end of its timestep, cheoa3 calls PSMILe prism_put routine to send Field11.

Finally, at the end of the run, cheoa3 performs the PSMILe finalisation call.

TOYOASIS3 coupling algorithm

The coupling algorithm between the TOYOASIS3 component models oceoa3, atmoa3, and cheoa3 is described here.

Table 6.1 lists the coupling fields exchanged between those 3 model components, giving the symbolic name used in each component and indicating whether the model produces the field (src) or receives it (tgt).

Figure 6.1 illustrates the coupling algorithm between the 3 TOYOASIS3 toy models for *Field₁*, *Field₃*, *Field₄*, *Field₁₀*, and *Field₁₁*.

	oceoa3	atmoa3	cheoa3	restart
Field1	SOSSTSST (src)	SISUTESU (tgt)		fldo1.nc
Field2	SOICECOV (src)	SIICECOV (tgt)		fldo2.nc
Field3	SOALBEDO (tgt)			SOALBEDO.nc
Field4	SONSHLDO (tgt)	CONSFTOT (src)		flda1.nc
Field5		COSHFTOT (src)		
Field6	SOWAFLDO (tgt)	COWATFLU (src)		flda3.nc
Field7	SORUNOFF (tgt)	CORUNOFF (src)		flda4.nc
Field8	SOZOTAUX (tgt)	COZOTAUX (src)		flda5.nc
Field9	SOMETAUY (tgt)	COMETAUY (src)		flda6.nc
Field10		COSENHFL (src)	SOSENHFL (tgt)	flda7.nc
Field11		COTHSHSU (tgt)	SOTHSHSU (src)	flda8.nc

Table 6.1: Coupling and I/O fields of the TOYOASIS3 coupled model. The symbolic name used in each toy model is given and it is indicated whether the model produces the field (src) or receives it (tgt).

$Field_1$ is sent from oceoa3 component to atmoa3 component at the coupling frequency dtF_1 defined by the user in the configuring file *namcouple*. As interpolation is needed between oceoa3 and atmoa3 grids, this exchange must go through OASIS3 interpolation process. In the *namcouple*, $Field_1$ field status must therefore be *EXPORTED* and the interpolation must be defined. If the user wants the field to be also automatically written to files before being sent (below the prism_put), and after being received (below the prism_get), he can choose the field status *EXPOUT*. In oceoa3 and atmoa3 codes, the prism_put and prism_get routines are respectively called every timestep with an argument corresponding to the time at the beginning of the timestep. The lag of $Field_1$, defined as 4 hours (14400 seconds) in the *namcouple*, is automatically added to the prism_put time argument; the prism_put called at the oceoa3 timestep preceeding the coupling period therefore matches the prism_get called in atmoa3 at the coupling period.

At the beginning of the run (i.e. at time = 0), the oceoa3 prism_put for $Field_1$ is not activated (as a positive lag is defined for $Field_1$) and OASIS3 automatically read $Field_1$ in its coupling restart file, fldo1.nc, and sends it to atmoa3 component after interpolation.

The exchange of $Field_2$ from oceoa3 to atmoa3 and $Field_4$, $Field_6$, $Field_7$, $Field_8$ and $Field_9$ from atmoa3 to oceoa3 follow exactly the same logic as for $Field_1$.

$Field_3$ as a status INPUT in the *namcouple*. $Field_3$ will therefore not be exchanged between two models but will be read from a file automatically below the target model oceoa3 prism_get calls, at the user-defined frequency in the input file also specified in the *namcouple*, SOALBEDO.nc.

$Field_5$ as a status of OUTPUT in the *namcouple*. It will therefore be only automatically written to a file at the user-defined frequency, below the source model atmoa3 prism_put calls. The name of the file will be automatically composed of the field symbolic name (here COSHFTOT) and of the begin and end dates of the run.

$Field_{10}$ and $Field_{11}$ are exchanged respectively from atmoa3 to cheoa3 and from cheoa3 to atmoa3. The fields status chosen by the user for those fields in the *namcouple* should therefore be IGNORED (or IGNOUT if the user wants the fields also automatically be written to files below the prism_put and after the prism_get). At the beginning of the run (i.e. at time = 0), the oceoa3 prism_get called to receive those fields will automatically read the fields in their corresponding coupling restart files flda7.nc and flda8.nc.

6.2.2 Compiling and Running TOYOASIS3

The TOYOASIS3 compiling and running environment is available in oasis3/examples/toyoasis3. Subdirectory /data contains auxiliary grid data files (see 5.2) and coupling restart files (see 5.3). Subdirectory /input contains the file cf_name_table.txt (see 5.1) the configuration file (see chapter 3) used in the classic (not parallel) mode, *namcouple*, and the configuration files used in the IPSL parallel




Figure 6.1: Exchange algorithm between the 3 TOYOASIS3 component models for fields $Field_1$, $Field_3$, $Field_4$, $Field_{10}$, and $Field_{11}$.

mode, *namcouple_0* and *namcouple_1*.

To run TOYOASIS3, first compile OASIS3 and the 3 TOYOASIS3 component models (see section 6.1). Go in directory `oasis3/examples/toyoasis3/src` and type `make`. This will automatically compile OASIS main executable and PSMILe library, if not done before hand, and the three component models `atmoa3.x`, `cheoa3.x` and `oceoa3.x`, using the header file specified in `oasis3/util/make_dir/make.inc`.

The next step is to adapt the “User’s section” of the running script `run_toyoasis3` in subdirectory `/script` and to launch it. The script `run_toyoasis3` supports Linux PC, NEC SX, IBM Power4, CRAYX1, CRAYXD1 and CRAYXT platforms (see `arch` variable). If your platform is not supported, the script will have to be adapted.

6.3 Known problems when compiling or running OASIS3 on specific platforms

- Notes on porting to BlueGeneL (XXX to be detailed):
- Notes on porting to BlueGeneP (XXX to be detailed):

Appendix A

The grid types for the transformations

As described in section 4, the different transformations in OASIS3 support different types of grids. The characteristics of these grids are detailed here.

1. Grids supported for the SCRIPR interpolations

- `'LR'` grid: The longitudes and the latitudes of 2D Logically-Rectangular (LR) grid points can be described by two arrays `longitude(i,j)` and `latitude(i,j)`, where `i` and `j` are respectively the first and second index dimensions. Stretched or/and rotated grids are LR grids. Note that A, B, G, L, Y, or Z grids are all particular cases of LR grids.
- `'U'` grid: Unstructured (U) grids do have any particular structure. The longitudes and the latitudes of 2D Unstructured grid points must be described by two arrays `longitude(nbr_pts,1)` and `latitude(nbr_pts,1)`, where `nbr_pts` is the total grid size.
- `'D'` grid The Reduced (D) grid is composed of a certain number of latitude circles, each one being divided into a varying number of longitudinal segments. In OASIS3, the grid data (longitudes, latitudes, etc.) must be described by arrays dimensioned `(nbr_pts,1)`, where `nbr_pts` is the total number of grid points. There is no overlap of the grid, and no grid point at the equator nor at the poles. There are grid points on the Greenwich meridian.

Appendix B

Changes between versions

Here is a list of changes between the different official OASIS3 versions.

B.1 Changes between `oasis3_3` and `oasis3-mct`

Much of the underlying implementation was refactored to leverage the Model Coupling Toolkit (MCT). All communication is now parallel and direct between models. The central oasis coupler running on it's own processors no longer plays a role. MCT is used to carry out interpolation in parallel on either the source or destination processors before or after communication. This location of the interpolation is set by the user as is the parallelization strategy. Several *namcouple* features have been deprecated in this version, but this version supports a *namcouple* format that mostly backwards compatible with OASIS3. The ability to read mapping files is now provided with the *namcouple* transformation MAPPING. Several other changes are highlighted below.

- General Features:
 - MPI2 job launching is NOT supported. Only MPI1 is allowed.
 - All coupling is now direct. Gets are blocking, puts are non-blocking. Get and put calls between models must match adequately to avoid deadlocks.
 - Like `oasis3`, the first coupling occurs at time=0 and the final coupling occurs at $(ncouplings-1)*dt$.
 - All input and output files must be netcdf, binary not supported
 - Ability to send and receive more than 1 field with a coupling operation. Use colon delimited field names in *namcouple* input.
 - A model output field (put) can be associated with more than one destination model (get). In that case, the source model only needs to put the data once and that data will arrive at multiple destinations as specified by *namcouple* input file. Different coupling frequencies and transforms are allowed for different coupling interactions of the same field. In addition, a single coupling field can be sent to the same destination model using multiple transforms as long as the destination field name is unique. For example
 - * TSURF from model 1 to TSURF in model 2
 - * TSURF from model 1 to TSURFAVG in model2
 - * TSURF from model 1 to TSURF in model 3

TSURF only needs to be put from model 1 “once” per timestep and the data will arrive at the destination at the coupling frequency and with transforms specified in the *namcouple* file for each coupling interaction. The inverse of this coupling feature is not allowed, a single destination (get) field CANNOT be associated with multiple source (put) fields.

- Only first order mapping methods are supported. Second order conserv and bicubic mapping is not supported.
- Vector mapping is currently not supported.
- Grid corners will NOT be compute automatically for LR grids if they are needed but not provided.
- Coupling Interfaces:
 - use mod_prism OR use mod_oasis instead of many use statements
 - * mod_prism retains prior interface names
 - * mod_oasis has updated interface names, changing “prism” to “oasis” and dropping “_proto” in the interface name
 - NO MPI2 support yet
 - prism_put_inquire not supported yet
 - prism_put_restart_proto not supported yet
 - prism_get_freq not supported yet
- *namcouple* Changes:
 - *namcouple* should be mostly backwards compatible, but several things are NOT used or supported and are ignored
 - * INIDATE is not used yet
 - * CALTYPE is not used
 - * AUXILARY is not used
 - * ONCE time transformation not supported
 - * REDGLO not supported
 - * INVERT not supported
 - * MASK not supported
 - * EXTRAP not supported
 - * CORRECT not supported
 - * INTERP not supported
 - * MOZAIC not supported
 - * FILLING not supported
 - * SUBGRID not supported
 - * MASKP not supported
 - * REVERSE not supported
 - * GLORED not supported
 - BLASOLD and BLASNEW only supports multiplication and addition terms to the fields, not field combinations.
 - IGNORED and INGOUT no longer needed, redundant with EXPORTED and EXPOUT use of IGNORED and INGOUT converts to EXPORTED and EXPOUT respectively
 - NLOGPRT sets debugging output levels for the coupling software and will lead to output to all log files.
 - Non-INSTANT LOCTRANS options will use a restart file
 - Fields of type OUTPUT now support a restart file argument to be compatible with LOC-TRANS restart files
 - SEQ setting has a slightly different meaning. SEQ is no longer needed to ensure correct coupling sequencing. Sequencing is dependent only on the order of get and put calls in models.

SEQ is never required. Use of SEQ allows the coupling layer to detect potential deadlocks before they happen and to exit gracefully. Use of SEQ is recommended when it makes sense to do so.

- There is a new MAPPING transform that allows use of previously generated mapping file for interpolation to be read in.
- Support to couple multiple fields via a single communication. This is supported through colon delimited field lists in the *namcouple* input, for example ATMTAUX:ATMTAUY:ATMHFLUX:ATMPREC TAUX:TAUY:HEATFLUX:FWFLUX in a single *namcouple* definition. All fields will use the *namcouple* settings for that definition. The *prism.get* and *prism.put* interfaces still only support sending (put) or receiving (get) one field at a time. Inside oasis, the fields are stored and a single mapping and send or receive is executed for all fields. This is useful in cases where multiple fields have the same coupling transforms or to reduce communication costs by aggregating multiple fields into a single communication.
- SCRIPR map generation supports a subset of options including BILINEAR, first order CONSERV, DISTWGT, and GAUSSWGT. No second order mapping is supported.

B.2 Changes between oasis3_3 and oasis3_prism_2_5

The changes between version oasis3_3 and version oasis3_prism_2_5 delivered in September 2006 are the following:

- Bug corrections:

- In `oasis3/lib/scip/src/remap_bilinear.f`, `remap_bicubic.f`, `remap_bilinear_reduced.f`, `remap_bicubic_reduced.F90`: (r2084) we observed a wrong behaviour of routines `remap_bilinear.f` and `remap_bicubic.f` on the NEC SX9 when compiled with NEC SX compiler revision 400. We got round this problem by adding an explicit instruction to prevent the vectorisation of one loop.
As `remap_bilinear_reduced.f` and `remap_bicubic_reduced.F90` have a very similar loop, we introduced the same instruction in these routines, although nothing specific was observed with them.
- `oasis3/lib/scip/src/remap_bicubic_reduced.F90`: (r1883) Two bugfixes: 1) Memory fault when a target point was falling on the before-last latitude circle or in the before-last latitude band of the source reduced grid; 2) Wrong neighbours for target points south of the before-last latitude circle (i.e in the last latitude band or Southern).
- `oasis3/lib/psmile/src/mod_psmile_io.F90`: (r2380) correction to ensure that when INVERT is used, the corner latitudes and longitudes are also inverted (and not only the center latitudes and longitudes as before).
- `oasis3/lib/scip/src/sciprmp.F`: (r1547) the calculation of the average for the line of points at the pole was bugged and did not have any effect. It is now debugged but commented.
- In `oasis3/lib/scip/src/vector.F90`:
 - correction of wrong sequences in declarations, at least for Intel & NAG compilers (thanks to L. Kornbluh from MPI); bug fix announced to the mailing list `diff-oasis@cerfacs.fr` on 31/10/2006.
 - (r1698 - 2008-08-20) bugfix to make sure that OASIS does not automatically calculates corners of target grid as this calculation is correct only for LR grids and target grid type is not known (thanks to S. Calmanti from Météo-France)
 - Modifications so that last 4 arguments of call to `grid_init` are always arrays even when corners are not defined (error detected with Intel Fortran V10.1.012 by Mike Rezny, SGI, Australia)

- In `oasis3/lib/clim.GSIP/src/CLIM.Init.Oasis.F`, correction of a wrongly positioned `#endif` (thanks to L. Kornblueh from MPI); bug fix announced to the mailing list `diff-oasis@cerfacs.fr` on 31/10/2006.
 - In `oasis3/lib/psmile/src/prism_enddef_proto.F`, the call to `MPI_Errhandler_set` was moved after the test on the value of `mpi_err` returned by `MPI_Buffer_Detach` (thanks to I. Bethke from NERSC); bug fix announced to the mailing list `diff-oasis@cerfacs.fr` on 14/12/2006.
 - Routine `oasis3/lib/psmile/src/prism_terminate_proto.F90` was modified to ensure proper deallocation of all allocated arrays (thanks to Adam Ralph from ICHEC)
 - In `oasis3/lib/scrip/remap_conserv.F`, small bugfix having no impact on the results, it just avoids misleading messages of type "Error sum wts map1:grid2_add ..." to be printed in the `cplout` log file.
 - In `oasis3/src/getfld.F`, `givfld.F`, `driver.f` and `closerst.F`, correction of a bug observed by A. Caubel from CEA for coupling fields having a sequence index greater than 1. For first iteration, closing of netcdf restart files is done in `driver.f` by calling new routine `closerst.F`. Closing is therefore removed from `getfld.F`. A minor correction (useless opening and closing of first netcdf restart file) was also added to `givfld.F`. Bug fix announced to the mailing list `diff-oasis@cerfacs.fr` on 09/02/2006.
 - In `oasis3/src/inipar.F`, bugfix for `SEQMODE` greater than 9 (thanks to T. Silva, Oregon State U.) and to avoid array overbound.
 - In `oasis3/util/make_dir/TopMakefileOasis3`: typo error: `libmpp_io.a` instead of `libmppio.a` (thanks to J.M. Epitalon from CERFACS)
 - In `oasis3/lib/psmile/src/prism_init_comp_proto.F`: initialisation of `iprcou` (bug fix thanks to J.M. Epitalon).
 - In `oasis3/src/filling.f`: rewind of file in the "AN" case (thanks to S. Calmanti from Météo-France)
 - In `oasis3/lib/psmile/src/mod_prismput_proto.F90`: bug fix to avoid array overbounds (bug identified T. van Noije from KNMI).
 - In `oasis3/lib/psmile/src/mod_psmile_io.F90`: thanks to A. Caubel from CEA, modification so that the grid point longitudes and latitudes do not have to appear before the corner longitudes and latitudes in the `grids.nc` file (revision 13/01/2009).
 - Bugfix in `oasis3/examples/testNONE/PROG/calc_errorfield.f90` to use the absolute value of the error in the non masked point mean error calculation.
 - Modifications in `oasis3/lib/clim/src/CLIM.Init.Oasis.F` and `oasis3/lib/psmile/src/prism_init_comp_proto.F` to allow communication log file (*.prt* files) for OASIS and/or component models to run on up to 9999 processes.
- Other major modifications
 - New directory structure. See section 1.2 for details.
 - Modification of many routines to allow using more than one OASIS3 executable in a coupled model resulting in pseudo-parallelisation of OASIS3 on a field-per-field basis. See section ?? for more detail.
 - New directory structure and update of compiling environnement to work with the new directory structure. In particular, the location of directory created for compilation (see `ARCHDIR` in the Makefile headers `make.xxx` in `oasis3/util/make_dir`) can be arbitrarily chosen by the user.
 - Modification of routine `oasis3/lib/scrip/src/remap_distwgt.F` so that `SCRIPR/DISTWGT` that has by default the same behaviour than `SCRIPR/BILINEAR`, `/BICUBIC` and `/CONSERV` (with `FRACNNEI` option) i.e. the non-masked nearest neighbour is used

for target grid points having their N nearest neighbour all masked. To reproduce the previous default behaviour, one has to compile with CPP key `NOT_NNEIGHBOUR`. See section 4.3 for details.

- Inclusion of an additional number below the `$NFIELDS` keyword in the *namcouple* (after the total number of fields exchanged, on the same line). This number, corresponding to the maximum number of `prism_def_var_proto` called by ANY component model in the coupled system, is needed only if it is greater than twice the number of fields listed in the *namcouple*; this may be the case if OASIS3 is used in pseudo-parallel mode or if fields declared with `prism_def_var_proto` call (and corresponding to sending - `prism_put_proto` call- or receiving - `prism_get_proto` call - actions in the component models) do not appear in the *namcouple* (in this case, the sending and receiving calls simply return without any action performed).
 - added options `GLBPOS`, `BASPOS`, `BASBAL` for the cooking stage transformation `CONSERV`. For details, see section 4.4.
 - New CPP key `TREAT_OVERLAY` : to ensure that if two cells of the source grid overlay, at least the one with the greater numerical index is masked (they also can be both masked). For example, if the grid line with `i=1` overlaps the grid line with `i=imax`, it is the latter that must be masked. When this is not the case with the mask defined in *masks.nc*, this CPP key ensures that these rules are respected. This is mandatory for `SCRIPR/CONSERV` remapping, see section 4.3.
 - Optimisation of `SCRIPR` (XXX to be detailed) interpolation weights-and-address files, thanks to CMCC. The weights and addresses are now read once per run and stored.
 - mode stats consommation Eric M (XXX to be detailed)
 - `mod_prism_get_comm` (XXX to be detailed)
 - Release of a new toy coupled model `TOYOASIS3`. This new toy model is described in 6.2. The toy model sources are available in `oasis3/src/mod/ocea3/`, `atmoa3`, `cheoa3`. Its running environment is available in `oasis3/src/mod/oasis3/examples/toyoasis3`. The grids of the `TOYOASIS3` component models and the interpolation performed have been updated compared to the previous toy coupled model `TOYCLIM` (which is no longer distributed with the official release).
 - In `oasis3/src/mod/oasis3/src/extrap.F`: optimisation (thanks to T. Schoenemeyer from NEC) and `idoitold` changed from 1000000 to 10000000 to support higher resolution grids (thanks to E.Maisonnavé from CERFACS). XXX and CMCC
 - In `oasis3/lib/psmile/src/mod_prism_grids_writing.F90`:
 - routine `prism_write_angle` was added to allow a component model to write the angle of its grid (see section 2.3 for more detail).
 - changed logical `netcdf` for `l_netcdf` (to run on NEC SX9, thanks to E. Maisonnavé, CERFACS)
 - In may 2007, we moved from CVS to SVN for source management.
- Other modifications
 - The names of the log files for the communication information **.prt** are now always ending with 4 digits indicating the rank of the component process (e.g. `modell.prt0002` or `modell.prt9999`) or the rank of the oasis process (e.g. `Oasis.prt0001` or `Oasis.prt0010`). These modifications impacted routines `oasis3/lib/clim/src/CLIM_Init_Oasis.F` and `oasis3/lib/psmile/src/prism_init_comp_proto.F`.
 - The following routines were modified for the NAG compiler: `oasis3/src/getfld.F`, `inipar.F`, `interp.F` and `oasis3/lib/scrip/src/remap_write.F`

- Some routines in `oasis3/lib/mpp_io/src/` were modified so that all debug and log messages are now written to stdout unit and to include modifications done in the OASIS4 version for CRAY pointers and for bundles.
- Added CPP key `__SILENT` CPP key to reduce log outputs to .prt files during the psmile library exchanges (thanks to S. Lorenz from MPI)
- In `oasis3/src/chkfld.f`, modified misleading comment about masked points written to *cplout* log file (thanks to T. Craig from BOM).
- In `oasis3/lib/psmile/src/mpp_io/src/mpp_io_mod.oa.F90` modified "lower-case" function so to avoid using "transfer" function (thanks to Mike Rezny for SGI Melbourne) which causes problem with pgf90 5.2.4, 6.1.3 or 7.0 in 64 bit mode, and with gfortran 4.2.1 and 4.2.5 SUSE Linux.

B.3 Changes between oasis3_prism_2_5 and oasis3_prism_2_4

The changes between version `oasis3_prism_2_5` and version `oasis3_prism_2_4` delivered in December 2004 are listed here after. Please note that those modifications should not bring any difference in the interpolation results, except for SCRIPR/DISTWGT (see below).

- Bug corrections:

- In `prism/src/lib/scrip/src/scriprmp.F`: initialisation of `dst_array(:)`; bug fix announced to the mailing list `diff-oasis@cerfacs.fr` on 02/02/2006.
- In `prism/src/lib/psmile/src/prism_enddef.proto.F` and `prism/src/lib/clim/src/CLIM_Start_MPI.F`: the call to `MPI_barrier` (that created a deadlock when not all processes of a component model were exchanging data with the coupler) was changed for a call to `MPI_wait` on the previous `MPI_Isend`; bug fix announced to the mailing list `diff-oasis@cerfacs.fr` on 02/23/2006.
- For SCRIPR/DISTWGT, in `prism/src/lib/scrip/src/remap_distwgt.f`: line 190 was repeated without epsilon modification; bug fix announced to the mailing list `diff-oasis@cerfacs.fr` on 03/21/2006.
- In `prism/src/lib/psmile/src/mod_prism_put_proto.F90`, for `prism_put_proto_r28` and `prism_put_proto_r24`, the reshape of the 2d field was moved after the test checking if the field is defined in the *namcouple* (thanks to Arnaud Caubel from LSCE).

- Modification in SCRIP interpolations

- For SCRIPR interpolations (see section 4.3), the value `1.0E+20` is assigned to target grid points for which no value has been calculated if `prism/src/lib/scrip/src/scriprmp.f` or `vector.F90` (for vector interpolation) are compiled with `ll_weightot = .true..`
- For SCRIPR/GAUSWGT: if routine `prism/src/lib/scrip/src/remap_gauswgt.f` is compiled with `ll_nnei=.true..`, the non-masked nearest neighbour is used for target point if all original neighbours are masked (see section 4.3).
- For SCRIPR/BICUBIC (routine `prism/src/lib/scrip/src/remap_bicubic.f`), the convergence criteria was modified so to ensure convergence even in single precision.
- For SCRIPR/CONSERV (routine `prism/src/lib/scrip/src/remap_conserv.f`), a test was added for non-convex cell so that integration does not stall.
- The routine `prism/src/lib/scrip/src/corners.F` was modified so to abort if it is called for the target grid, as the automatic calculation of corners works only for Logically-Rectangular (LR) grids and as the target grid type is unknown. If needed, the reverse remapping, in which the current target grid become the source grid, can be done .

- Other important modifications

- A new PSMILE routine `prism/src/lib/psmile/src/prism_get_freq.F` was added; this routine can be used to retrieve the coupling period of field (see section ??).
- The routines of the `mpp_io` library in `prism/src/lib/mpp_io` changed name and were merged with the `OASIS4 mpp_io` library.
- Routine `prism/src/mod/oasis3/src/extrap.F` was modified to ensure that the extrapolation works even if the `MASK` value is very big (thanks to J.M. Epitalon).
- In the *namcouple*, there is no need anymore to define a lag greater than 0 (e.g. `LAG=+1`) for fields in mode `NONE`.
- Diverse modifications were included for successful compilation with NAGW compiler: non portable use of “kind”, etc. (thanks to Luis Kornblueh from MPI).
- In `prism/src/lib/psmile/mod-prism_get_proto.F90`, a potential deadlock was removed (the master process was sending a message to itself)(thanks to Luis Kornblueh from MPI).
- Routine `prism/src/lib/scrip/src/scriprmp_vector.F90` was completely rewritten for more clarity.
- Obsolete transformations `INVERT` and `REVERSE` were removed from the toy coupled model `TOYCLIM` (in file `prism/util/running/toyclim/input/namcouple`. This change does not affect the statistics printed in the `cplout` but changes the orientation of some fields in the NetCDF output files (see the results in `prism/data/toyclim/outdata`).
- Other minor modifications:
 - In `prism/src/lib/psmile/src/prism_enddef_proto.F`, allocation is done only for `rg_field_trans` or `dg_field_trans` depending on precision for `REAL` (but not for both, to save memory).
 - In few routines in `prism/src/lib/clim` and in `prism/src/mod/oasis3`, parentheses were added to make sure that `&&` has priority over `||` in CPP instructions (thanks to A. Caubel from LSCE).
 - Routines `scrip/src/corners.f`, `netcdf.f`, and `scriprmp.f` were renamed `corners.F`, `netcdf.F`, `scriprmp.F` and the line “`INCLUDE 'netcdf.inc'`” was changed for “`#include <netcdf.inc>`”

B.4 Changes between `oasis3_prism_2_4` and `oasis3_prism_2_3`

The changes between versions tagged `oasis3_prism_2_4` and `oasis3_prism_2_3` delivered in July 2004 are the following:

- Update of compiling and running environments with version `prism_2-4` of PRISM Standard Compiling Environment (SCE) and PRISM Standard Running Environment (SRE), which among other improvements include the environments to compile and run on the CRAY X1 (see the directories with `<node>=baltic1`), thanks to Charles Henriet from CRAY France, and on a Linux station from Recherche en Prévision Numérique (Environnement Canada, Dorval, Canada) (see the directories with `<node>=armc28`).
- `prism/src/mod/oasis3/src/iniiof.F`: the opening of the coupling restart files is done only if the corresponding field has a lag greater than 0; note that this implies that all fields in mode `NONE` must now have a lag greater than 0 (e.g. `LAG=+1`) (thanks to Veronika Gayler from M&D).
- `prism/src/lib/psmile/src/prism_def_var_proto.F`: contrary to what was previously described in the documentation, `PRISM_Double` is not supported as 7th argument to describe the field type; `PRISM_Real` must be given for single or double precision real arrays.

- `prism/src/mod/oasis3/src/inipar.F90`: For upward compatibility of SCRIPR interpolation, “VECTOR” is still accepted in the *namcouple* as the field type and leads to the same behaviour as before (i.e. each vector component is treated as an independent scalar field). To have a real vector treatment, one has to indicate “VECTOR_I” or “VECTOR_J” (see section 4.3).
- Bug corrections in:
 - `prism/src/lib/scrip/src/scriprmp_vector.F90`: In some cases, some local variables were not deallocated and variable `dimid` was declared twice.
 - `prism/src/lib/psmile/src/mod_psmile_io.F90`: correct allocation of array hosting the longitudes (thanks to Reiner Vogelsang from SGI Germany).
 - `prism/src/lib/psmile/src/write_file.F90`: to remove a deadlock on some architecture (thanks to Luis Kornblueh from MPI).
 - `prism/src/lib/psmile/src/prism_enddef_proto.F`: the error handler is now explicitly set to `MPI_ERRORS_RETURN` before the call to `MPI_Buffer_Detach` to avoid abort on some architecture when the component model is not previously attached to any buffer (thanks to Luis Kornblueh from MPI).
 - `prism/src/lib/scrip/src/remap_conserv.f` (thanks to Veronika Gayler from M&D).
 - `prism/src/mod/oasis3/src/inicmc.F`
 - `prism/src/lib/scrip/src/remap_distwgt.f`

B.5 Changes between `oasis3_prism_2_3` and `oasis3_prism_2_2`

The changes between versions tagged `oasis3_prism_2_3` delivered in July 2004 and `oasis3_prism_2_2` delivered in June 2004 are the following:

- Bug correction of the previous bug fix regarding ordering of grid and data information contained in I/O files when `INVERT` or `REVERSE` transformations are used: the re-ordering now occurs only for source field if `INVERT` is used, and only for target field if `REVERSE` is used.
- LGPL license: OASIS3 is now officially released under a Lesser GNU General Public License (LGPL) as published by the Free Software Foundation (see `prism/src/mod/oasis3/COPYRIGHT` and `prism/src/mod/oasis3/src/couple.f`)
- Upgrade of compiling and running environments: The compiling and running environments have been upgraded to the PRISM Standard Compiling and Running Environment version dated August 5th 2004, that should be very close to “`prism_2-3`”.
- Treatment of vector fields: The interpolation algorithms using the SCRIP library now support vector fields, including automatic rotation from local to geographic coordinate system, projection in Cartesian coordinate system and interpolation of 3 Cartesian components, and support of vector components given on different grids. New routines have been added in `prism/src/lib/scrip/src/scriprmp_vector.F90` and `rotations.F90`. For more detail, see SCRIPR in section 4.3.
- All include of `mpif.h` are now written ‘`#include <mpif.h>`’.
- The output format of `CHECKIN` and `CHECKOUT` results is now E22.7

B.6 Changes between `oasis3_prism_2_2` and `oasis3_prism_2_1`

The changes between versions tagged `oasis3_prism_2_2` delivered in June 2004 and `oasis3_prism_2_1` delivered to PRISM in April 2004 are the following:

- Bug corrections
 - `INTERP/GAUSSIAN` and `SCRIPR/GAUSWGT` transformations work for ‘U’ grids.

- The grid and data information contained in I/O files output by the PSMILe library have now a coherent ordering even if INVERT or REVERSE transformations are used.
- OASIS3 and the TOYCLIM coupled model are ported to IBM Power4 and Linux Opteron, which are now included in the Standard Compiling and Running Environments (SCE and SRE).
- SIPC technique communication is re-validated.
- `Clim_MaxSegments = 338` in `prism/src/lib/clin/src/mod_clim.F90` and in `prism/src/lib/psm`. 338 is presently the largest value needed by a PRISM model.
- `MPI_BSend`: below the call to `prism_enddef_proto`, the PSMILe tests whether or not the model has already attached to an MPI buffer. If it is the case, the PSMILe detaches from the buffer, adds the size of the pre-attached buffer to the size needed for the coupling exchanges, and reattaches to an MPI buffer. The model own call to `MPI_Buffer_Attach` must therefore be done before the call to `prism_enddef_proto`. Furthermore, the model is not allowed to call `MPI_BSend` after the call to `prism_terminate_proto`, as the PSMILe definitively detaches from the MPI buffer in this routine. See the example in the toyatm model in `prism/src/mod/toyatm/src`.

B.7 Changes between oasis3_prism_2.1 and oasis3_prism_1.2

The changes between versions tagged `oasis3_prism_1.2` delivered in September 2003 and `oasis3_prism_2.1` delivered to PRISM in April 2004 are the following:

- Bug corrections
 - Thanks to Eric Maisonnave, a bug was found and corrected in `/prism/src/lib/scrip/src/scriprmp.f`: “sou_mask” and “tgt_mask” were not properly initialised if weights and addresses were not calculated but read from file.
 - Some deallocation were missing in `prism_terminate_proto.F` (“ig_def_part”, “ig_length_part”, “cg_ignout_field”).
 - Thanks to Arnaud Caubel, a bug was found and corrected in `/prism/src/lib/psmile/src/write_file.F90`. In case of parallel communication between a model and OASIS3 main process, the binary coupling restart files were not written properly (NetCDF coupling restart files are OK).

- Routines renamed

The routines `preproc.f`, `extrap.f`, `iniiof.f` in `prism/src/mod/oasis3/src` were renamed to `preproc.F`, `extrap.F`, `iniiof.F`, as a CPP key ‘key_openmp’ was added. Please note that this key, allowing openMP parallelisation, is not fully tested yet.

- Modifications in the *namcouple*

- The third entry on the field first line now corresponds to an index in the new auxiliary file *cf_name_table.txt* (see sections 3.3 and 5.1).
- For IGNORED, IGNOUT and OUTPUT fields, the source and target grid locator prefixes must now be given on the field second line (see section ??)

- A new auxiliary file *cf_name_table.txt*

For each field, the CF standard name used in the OASIS3 log file, *cplout*, is now defined in an additional auxiliary file *cf_name_table.txt* not in `inipar.F` anymore. This auxiliary file must be copied to the working directory at the beginning of the run. The user may edit and modify this file at her own risk. In *cf_name_table.txt*, an index is given for each field standard name and associated units. The appropriate index has to be indicated for each field in the *namcouple* (third entry on the field first line, see section 3.3).

This standard name and the associated units are also used to define the field attributes “long_name” and “units” in the NetCDF output files written by the PSMILe for fields with status EXPOUT, IGNOUT and OUTPUT.

For more details on this auxiliary file, see section 5.1.

- Many timesteps for mode NONE

In mode NONE, OASIS3 can now interpolate at once all time occurrences of a field contained in an input NetCDF file. The time variable in the input file is recognized by its attribute “units”. The acceptable units for time are listed in the `udunits.dat` file¹. This follows the CF convention.

The keyword `$RUNTIME` in the *namcouple* has to be the number of time occurrences of the field to interpolate from the input file. The “coupling” period of the field (4th entry on the field first line) must be always “1”. Note that if `$RUNTIME` is smaller than the total number of time occurrences in the input file, the first `$RUNTIME` occurrences will be interpolated.

For more details, see section ??.

- Model grid data file writing

The grid data files *grids.nc*, *masks.nc* and *areas.nc* can now be written directly at run time by the component models, if they call the new routines `prism_start_grids_writing`, `prism_write_grid`, `prism_write_corner`, `prism_write_mask`, `prism_write_area`, `prism_terminate_grids_writing`.

The writing of those grid files by the models is driven by the coupler. It first checks whether the binary file *grids* or the netCDF file *grids.nc* exists (in that case, it is assumed that *areas* or *areas.nc* and *masks* or *masks.nc* files exist too) or if writing is needed. If *grids* or *grids.nc* exists, it must contain all grid information from all models; if it does not exist, each model must write its grid informations in the grid data files.

See section 2.3 for more details.

- Output of CF compliant files

The NetCDF output files written by the PSMILE for fields with status `EXPOUT`, `IGNOUT` and `OUTPUT` are now fully CF compliant.

In the NetCDF file, the field attributes “long_name” and “units” are the ones corresponding to the field index in *cf_name_table.txt* (see above and section 5.1). The field index must be given by the user as the third entry on the field first line in the *namcouple*.

Also, the latitudes and the longitudes of the fields are now automatically read from the grid auxiliary data file *grids.nc* and written to the output files. If the latitudes and the longitudes of the mesh corners are present in *grids.nc*, they are also written to the output files as associated “bounds” variable. This works whether the *grids.nc* is given initially by the user or written at run time by the component models (see above). However, this does not work if the user gives the grid definition in a binary file *grids*.

- Removal of pre-compiling key “key_BSend”

The pre-compiling key “key_BSend” has been removed. The default has changed: by default, the buffered MPI.BSend is used, unless `NOBSEND` is specified in the *namcouple* after `MPI1` or `MPI2`, in which case the standard blocking send `MPI_Send` is used to send the coupling fields.

¹See <http://www.unidata.ucar.edu/packages/udunits/udunits.dat>

Appendix C

Coupled models realized with OASIS

Table C.1: Use of OASIS3: centres, coupled models and computing platforms

Centre	Coupled model	Platform
CERFACS (FR)	ARPEGE4_T63/NEMO-ORCA2-LIM/TRIP	Linux Cluster, CRAY XD1, VPP5000, NEC SX6-SX8
CERFACS (FR)	ARPEGE_T63/NEMIX-ORCA2	CRAY XD1
CERFACS (FR)	ARPEGE_T359/NEMO-ORCA012	NEC SX9
CERFACS (FR)	ARPEGE_T799/NEMIX-ORCA025	Bullx, Altix ICE
Météo-France (FR)	ARPEGE_V4.6/NEMO-ORCA2/NEMOmed8	NEC-SX8
Météo-France (FR)	ALADIN-Climat/NEMOmed8/TRIP	NEC-SX9
Météo-France (FR)	ARPEGE_V5/NEMOV3-ORCA1/TRIP	NEC-SX9
Météo-France (FR)	ARPEGE_V5.1/NEMO1	IBM Power 6
IPSL (FR)	LMDZ/NEMO-ORCA2/ - LMDZmed/NEMOmed8	
IPSL (FR)	LMDz(144x142)/NEMO-ORCA2	NEC SX8
OMP (FR)	MESO-NH/SYMPHONIE	Linux Opteron cluster
LGGE (FR)	MAR/NEMO-LIM	
ECMWF	IFS_T399/NEMO-ORCA1	IBM Power 6
MPI-M (DE)	ECHAM5/MPIOM	IBM Power 6, SUN Linux
MPI-M (DE)	REMO/MPIOM	IBM Power 575
Met Office (UK)	UM Atm(192x145)/NEMO-ORCA1	IBM Power6, NEC SX6
Met Office (UK)	UM Atm(432x325)/NEMO-ORCA025	IBM Power6
NCAS/Reading (UK)	ECHAM4/NEMO-ORCA2	NEC SX6-SX8
NCAS/Reading (UK)	HadAM3/NEMO-ORCA2	NEC SX8
IFM-GEOMAR (DE)	ECHAM5_T63/NEMO-ORCA2	NEC SX6, SX8, SX9
CMCC (IT)	ECHAM5_T31L31/OPA8.2-ORCA2	NEC SX9
CMCC (IT)	ECHAM5_T159L31/OPA8.2-ORCA2	IBM Power6
CMCC (IT)	ECHAM5_T63L95/OPA8.2-ORCA2	
CMCC (IT)	ECHAM5_T159/OPA8.2-ORCA2/ - NEMOMed_1/16	
ENEA (IT)	RegCM/MITgcm	IBM-SP5
LMNCP (IT)	WRF/ROMS	
SMHI (SE)	EC-Earth: IFS_T159/NEMO-ORCA1	Linux Cluster
SMHI (SE)	EC-Earth: IFS_T799/NEMO-ORCA025	Ekman AMD Opteron
KNMI (NL)	ECHAM5/MPIOM	NEC SX-8
KNMI (NL)	EC-Earth: IFS_T159/NEMO-ORCA2	SGI Altix, IBM Power
DMI (DK)	ECHAM(global)/HIRLAM (reg)	NEC SX6
U.Bergen (NO)	MM5/ROMS	
ICHEC (IE)	EC-Earth: IFS_T159/NEMO-ORCA1	SGI Altix ICE, Bull clust.
ICHEC (IE)	ROMS/WRF	
NUI Galway (IE)	ECHAM5/REMO/MPI-OM	

Continued on next page

Table C.1 – continued from previous page

Centre	Coupled model	Platform
ETH (CH)	COSMO-CLM/CLM	IBM Power 6 - Intel Itanium Linux Fedora core 14 IBM Power5 Linux PC IBM IBM Regata Series Linux cluster Linux cluster NEC SX8 NEC SX8 IBM cluster SGI Linux_x64 SUN + SGI clusters NEC SX6, SUN NEC SX-6 SGI O3400 - Compaq5
ETH (CH)	COSMO-CLM/ROMS	
U. Castille (ES)	PROMES/U. Madrid ocean	
NHM Service (RS)	ECHAM5/MPIOM	
CMC (CA)	GEM/NEMO	
UQAM (CA)	GEMDM 3.3.2/NEMOv_2.3.0	
MM (MA)	ARPEGE-Climat_V5.1/NEMO-ORCA1	
INMT (TN)	ARPEGE-Climat_V5.1/NEMO-ORCA2	
Oregon St U (USA)	PUMA/UVic	
Hawaii U (USA)	ECHAM4/POP	
JAMSTEC (JP)	ECHAM/OPA8.2	
JAMSTEC (JP)	ECHAM5.T106L31/NEMO_0.5	
Met.Nat.Center (CN)	GRAPES (201x161)/ECOM-si	
IAP (CN)	CREM(reg.)/POM2000	
IAP (CN)	ECHAM/MPIOM	
CSIRO (AU)	ACCESS : UMv7.3/MOM4p1/CICE	
BoM (AU)	BAM3/ACOM2	
BoM (AU)	TCLAPS/MOM	
U Tasmania (AU)	Data atm. model/MOM4	

Table C.2: Use of OASIS2: centres, coupled models and computing platforms

Centre	Coupled model	Platform
CERFACS (FR)	ARPEGE3 OPA8.2/LIM 2deg	VPP5000
CERFACS (FR)	ARPEGE3 OPA8.1	VPP700
CERFACS (FR)	ARPEGE2 OPAICE	CRAY C90
IPSL (FR)	LMDz 96x71x19 - OPA/ORCA2	VPP5000
IPSL (FR)	LMDz 72x45x19 OPA/ORCA4	VPP5000
IPSL (FR)	LMDZ 120X90 - OPA ATL3 1/3	
IPSL (FR)	LMDZ 120X90 - OPA ATL1 1 deg	
LODYC (FR)	IFS T195L31 OPA8.1	
LODYC (FR)	ECHAM4 T30/T42 L14 OPA/ORCA2	
Météo-France (FR)	ARPEGE medias OPAMED 1/8 deg	VPP5000
Météo-France (FR)	ARPEGE3 - OPA 8.1 + Gelato	VPP5000
Météo-France (FR)	ARPEGE 2 T31L19 - OPA8 TDH	CRAY J90
ECMWF	IFS T63/T255 - E-HOPE 2deg/1deg	IBM Power 4
ECMWF	IFS Cy23r4 T159L40 - E-HOPE 256L29	VPP700
ECMWF	IFS Cy23r4 T95L40 - E-HOPE 256L29	VPP700
ECMWF	IFS Cy15r8 T63L31 - E-HOPE 128L20	VPP300
MPI (DE)	ECHAM5 T42/L19 - C-HOPE T42+L20	NEC-SX
MPI (DE)	PUMAT42/L19 - C-HOPE 2deg GIN	NEC-SX
MPI (DE)	EMAD - E-HOPE T42+L20	CRAY C-90
MPI (DE)	ECHAM5 T42/L19 - E-HOPE T42+L20	NEC-SX
MPI (DE)	ECHAM4 T30/L19 - E-HOPE T42+L20	CRAY T90
MPI (DE)	ECHAM4 T30/L19 - E-HOPE T42+L20	CRAY C90
SMHI (SE)	RCA-HIRLAM (reg.) - RCO-OCCAM (reg.)	
CGAM (UK)	HadAM3 2.5x3.75 L20 - OPA/ORCA2	T3E
SOC (UK)	Interm. Atm. GCM - OCCAM-Lite	
NOC (UK)	FORTE : IGCM3 (128 x 64) - MOMA (180 x 88)	
JPL (USA)	TRIDENT QTCM	Linux workstation
IRI (USA)	ECHAM4 - MOM3	SGI Origin - IBM Power3
JAMSTEC (JP)	ECHAM4 - OPA 8.2	NEC SX5
BMRC (AU)	BAM3 - ACOM2	NEC SX6

Bibliography

- [Cassou et al 1998] Cassou, C., P. Noyret, E. Sevault, O. Thual, L. Terray, D. Beaucourt, and M. Imbard: Distributed Ocean-Atmosphere Modelling and Sensitivity to the Coupling Flux Precision: the CATHODE Project. *Monthly Weather Review*, 126, No 4: 1035-1053, 1998.
- [Guilyardi et al 1995] Guilyardi, E., G. Madec, L. Terray, M. Déqué, M. Pontaud, M. Imbard, D. Stephenson, M.-A. Filiberti, D. Cariolle, P. Delecluse, and O. Thual. Simulation couplée océan-atmosphère de la variabilité du climat. *C.R. Acad. Sci. Paris*, t. 320, série IIa:683–690, 1995.
- [Jacob et al 2005] Jacob, R., J. Larson, and E. Ong: MxN Communication and Parallel Interpolation in CCSM3 Using the Model Coupling Toolkit. *Int. J. High Perf. Comp. App.*, 19(3), 293-307 2005
- [Jones 1999] Jones, P.: Conservative remapping: First- and second-order conservative remapping, *Mon Weather Rev*, 127, 2204-2210, 1999.
- [Larson et al 2005] Larson, J., R. Jacob, and E. Ong: The Model Coupling Toolkit: A New Fortran90 Toolkit for Building Multiphysics Parallel Coupled Models. *Int. J. High Perf. Comp. App.*, 19(3), 277-292, 2005
- [Noyret et al 1994] Noyret, P., E. Sevault, L. Terray and O. Thual. Ocean-atmosphere coupling. *Proceedings of the Fall Cray User Group (CUG) meeting*, 1994.
- [Pontaud et al 1995] Pontaud, M., L. Terray, E. Guilyardi, E. Sevault, D. B. Stephenson, and O. Thual. Coupled ocean-atmosphere modelling - computing and scientific aspects. In *2nd UNAM-CRAY supercomputing conference, Numerical simulations in the environmental and earth sciences* Mexico-city, Mexico, 1995.
- [Sevault et al 1995] Sevault, E., P. Noyret, and L. Terray. Clim 1.2 user guide and reference manual. *Technical Report TR/CGMC/95-47*, CERFACS, 1995.
- [Terray and Thual 1995b] Terray, L. and O. Thual. Oasis: le couplage océan-atmosphère. *La Météorologie*, 10:50–61, 1995.
- [Terray and Thual 1993] Terray, L. and O. Thual. Coupled ocean-atmosphere simulations. In *High Performance Computing in the Geosciences, proceedings of the Les Houches Workshop* F.X. Le Dimet Ed., Kluwer Academic Publishers B.V, 1993.
- [Terray et al 1995] Terray, L., E. Sevault, E. Guilyardi and O. Thual OASIS 2.0 Ocean Atmosphere Sea Ice Soil User's Guide and Reference Manual *Technical Report TR/CGMC/95-46*, CERFACS, 1995.
- [Terray et al 1995b] Terray, L. O. Thual, S. Belamari, M. Déqué, P. Dandin, C. Lévy, and P. Delecluse. Climatology and interannual variability simulated by the arpege-opa model. *Climate Dynamics*, 11:487–505, 1995
- [Terray et al 1999] Terray, L., S. Valcke and A. Piacentini: OASIS 2.3 Ocean Atmosphere Sea Ice Soil, User's Guide and Reference Manual, *Technical Report TR/CGMC/99-37*, CERFACS, Toulouse, France, 1999.
- [Valcke et al 2004] Valcke, S., A. Caubel, R. Vogelsang, and D. Declat: OASIS3 User's Guide (oasis3_prism_2-4), *PRISM Report No 2, 5th Ed.*, CERFACS, Toulouse, France, 2004.
- [Valcke et al 2003] Valcke, S., A. Caubel, D. Declat and L. Terray: OASIS3 Ocean Atmosphere Sea Ice Soil User's Guide, *Technical Report TR/CGMC/03-69*, CERFACS, Toulouse, France, 2003.
- [Valcke et al 2000] Valcke, S., L. Terray and A. Piacentini: OASIS 2.4 Ocean Atmosphere Sea Ice Soil, User's Guide and Reference Manual, *Technical Report TR/CGMC/00-10*, CERFACS, Toulouse, France, 2000.